

Rules of Thumb for Species Occupancy Models

Mark Logie & Tom August

12 October, 2018

Contents

Introduction	1
Metrics	2
Data quality metrics	2
Model quality metrics	2
Data	2
Read in CIRRUS outputs	2
Posteriors	6
Raw data files	8
Classification of Models	14
Create metrics	14
Consultation on model outputs	17
Classification tree	21
Classifying the data	24
Examining the classification variables	24
Correlation between variables	27
Input data requirement	28
Classification trees	30
Equally weighted decision tree	30
10:1 Good:bad tree - Aspirational Target	38
1:10 Good:bad tree - bare minimum to try running a model	42
How many species can we model?	44
For the UK	44
By habitat	50
By region	55
Conclusions	59

Built with 3.4.1

Introduction

The aim of this document is to start development for rules of thumb around occupancy models. Which species can we make reliable models for? How many records do we need? How large an area? how many visits? Having a greater understanding of how these variables effect the quality of the models produced will allow us to develop rules of thumb for which species can be modelled and at which scale. This will also help us to identify key data gaps.

Metrics

To develop these rules of thumb we need two sets of metrics.

Data quality metrics

The metrics describe the data that we have for a species and can be used to predict the quality of the model:

- The 50th/90th percentile of the number of records of the focal species per year [`median/P90`]
- The 50th/70th/80th/90th percentile of the number of visits to a site each year, for sites where the species has been observed in that year (i.e. including visits where the focal species was not recorded) [`visits_median/P70/P80/P90`]
- The proportion of years in which the focal species has been recorded [`prop_of_years`]
- The proportion of all site:year combinations, where the focal species is observed, that have > 1 visit [`prop_repeats_grp`]
- Considering all the lists where the focal species was recorded, the proportion of lists that had length 1 (i.e. records only of the focal species) [`prop_list_one`]
- Considering all visits for the ‘taxonomic’ group in the dataset, the proportion of all visits that did not record the focal species [`prop_abs`]
- Considering the visits for the ‘taxonomic’ group where the focal species was not observed, this is the proportion of visits with list length > 1 [`prop_abs_list`]*

*Note: this variable is excluded in the final decision trees. See ‘Classifying the data: Examining the classification variables’

Model quality metrics

These metrics measure the quality of the model produced and are what we will use to develop the rules of thumb:

- The precision of the trend where `Precision <- 1/(sd(long_term_gr)^2)` and `long_term_gr` is a vector of annual growth rates where the `length()` of `long_term_gr` is equal to the number of model iterations [`precision_growth_rate`]
- Average precision of the year estimates where precision is defined in the same way as above [`mean_year_precision`]
- Proportion of years with converged Rhat [`PropYrConverged`]

Data

The data are in a few different locations and in most cases we need to go to more than one data source to get both sets of metrics. There is also a difference in the location of outputs, some are on Charlie’s Cirrus folder, some in her JASMIN folder and some are on LOTUS still.

Read in CIRRUS outputs

First, set up the location of the files to be read in

```
files_are_here <-  
  "W:/PYWELL_SHARED/Pywell Projects/BRC/Charlie/1.c. New Model Rerun/4. Outputs/CIRRUS"  
  
rdata_files <- list.files(files_are_here,  
                           pattern = '.rdata$',  
                           recursive = TRUE,
```

```

        full.names = TRUE)

# The files we want are those in the 20000_update folder
rdata_files <- rdata_files[grep('20000_update', rdata_files)]

```

Count up the number of files and set up some variables

```

num_files <- length(rdata_files)
speciesName <- quant25 <- Spnsite <- Spnvisits <- nyears <-
  PropYrConverged <- Totalnvisits <- Totalnsites <- FirstYrConverged <-
  LastYrConverged <- c(rep(' ', num_files*2))

```

Get what data we can from the model outputs

```

plot_species <- function(new_data, x, species, main = ''){

  new_data <- as.data.frame(new_data)
  new_data[, 'year'] <-
    (Year = (x$min_year - 1) +
      as.numeric(gsub("psi.fs", "", gsub("\\\\[|\\\\]", "", row.names(new_data)))))

  # rename columns, otherwise ggplot doesn't work properly
  names(new_data) <- gsub("2.5%", "quant_025", names(new_data))
  names(new_data) <- gsub("97.5%", "quant_975", names(new_data))

  # Add rhat T/F column
  new_data$rhat_threshold[new_data$Rhat < 1.1] <- 'Good (<1.1)'
  new_data$rhat_threshold[new_data$Rhat > 1.1] <- 'Bad (>1.1)'

  # plot the yearly predicted proportion of occupied sites
  # plot with error bars based on 95CI
  ggplot(new_data, aes_string(x = "year", y = "mean")) +
    theme_bw() +
    geom_ribbon(data = new_data,
                aes_string(group = 1, ymin = "quant_025", ymax = "quant_975"),
                alpha = 0.2) +
    geom_line(size = 1, col = "black") +
    geom_point(size = 4, aes(col = rhat_threshold)) +
    scale_color_manual(name = 'Rhat',
                       values = c('Bad (>1.1)' = 'red', 'Good (<1.1)' = 'blue')) +
    ylab("Occupancy") +
    xlab("Year") +
    scale_y_continuous(limits = c(0, 1)) +
    ggtitle(main) +
    theme(plot.title = element_text(lineheight = .8, face = "bold"),
          legend.position = 'bottom')
  ggsave(filename = paste0('metrics/plots/', species, '.png'),
         width = 6.71, height = 4.64)
}

for(i in 1:num_files){
  # Load the file
  load(rdata_files[i])
  # Load the summary table
  bugsSummary <- out$BUGSoutput$summary
}

```

```

## ~~~ Useful information ~~~
speciesName[i] <- out$SPP_NAME

## ~~~ Explanatory variables ~~~
# Number of sites species is observed at
Spnsite[i] <- out$species_sites
# Number of visits on which the species is observed
Spnvisits[i] <- out$species_observations
# Number of years
nyears[i] <- length(out$max_year:out$min_year)
# Total visits
Totalnvisits[i] <- out$nvisits
# Total sites
Totalnsites[i] <- out$nsites
## Rhats ##
yearValues <- bugsSummary[grep1('^\$psi.fs\\\[', rownames(bugsSummary)), ]
# Proportion of years converged
PropYrConverged[i] <- sum(yearValues[, 'Rhat'] < 1.1) / nrow(yearValues)
# First/Last year converged?
FirstYrConverged[i] <- yearValues[, 'Rhat'][1] < 1.1
LastYrConverged[i] <- tail(yearValues[, 'Rhat'], 1) < 1.1

plot_species(yearValues, out, speciesName[i])

rm(list = c('out'))

cat(paste0('First pass - Finished processing file ', i, ' of ', num_files, '\n'))
}

# Now do just the last 10 years
for(i in 1:num_files){
  # Load the file
  load(rdata_files[i])

  # Load the summary table
  bugsSummary <- out$BUGSoutput$summary

  ## ~~~ Useful information ~~~
  speciesName[i+num_files] <- paste0(out$SPP_NAME, '_last10yrs')

  nyears[i+num_files] <- 10
  ## Rhats ##
  yearValues <- bugsSummary[grep1('^\$psi.fs\\\[', rownames(bugsSummary)), ]
  # subset to last 10 years
  yearValues <- tail(yearValues, 10)
  # Proportion of years converged
  PropYrConverged[i+num_files] <- sum(yearValues[, 'Rhat'] < 1.1) / nrow(yearValues)
  # First/Last year converged?
  FirstYrConverged[i+num_files] <- yearValues[, 'Rhat'][1] < 1.1
  LastYrConverged[i+num_files] <- tail(yearValues[, 'Rhat'], 1) < 1.1

  plot_species(yearValues, out, speciesName[i+num_files])
}

```

```

rm(list = c('out'))

cat(paste0('Second pass - Finished processing file ', i, ' of ', num_files, '\n'))
}

Spnsite   <- as.numeric(Spnsite)
Spnvisits <- as.numeric(Spnvisits)
quant25 <- as.numeric(quant25)
nyears <- as.numeric(nyears)
Totalnvisits <- as.numeric(Totalnvisits)
PropYrConverged <- as.numeric(PropYrConverged)

SpvisitPerSite <- Spnvisits/Spnsite
SpvisitPerSite[is.na(SpvisitPerSite)] <- 0

SpvisitPerYear <- Spnvisits/nyears
SpvisitPerYear[is.na(SpvisitPerYear)] <- 0

TotalvisitPerYear <- Totalnvisits/nyears
TotalvisitPerYear[is.na(TotalvisitPerYear)] <- 0

SpeciesData <- data.frame(speciesName, Spnsite, Spnvisits, SpvisitPerSite,
                           SpvisitPerYear, Totalnvisits, Totalnsites,
                           TotalvisitPerYear, nyears, FirstYrConverged,
                           LastYrConverged, PropYrConverged)

write.csv(SpeciesData,
          file = 'Results/metrics/model_data.csv', row.names = FALSE)

model_data <- read.csv('Results/metrics/model_data.csv')
str(model_data)

## 'data.frame': 13640 obs. of 12 variables:
## $ speciesName : Factor w/ 9266 levels "", "Lamprochromus bifasciatus", ... : 5461 5463 5465 5467 ...
## $ Spnsite      : int  51 336 16 852 237 75 11 0 355 6 ...
## $ Spnvisits    : int  94 785 35 1804 450 168 50 0 965 45 ...
## $ SpvisitPerSite : num  1.84 2.34 2.19 2.12 1.9 ...
## $ SpvisitPerYear : num  2 16.702 0.745 38.383 9.574 ...
## $ Totalnvisits : int  11260 11260 11260 11260 11260 11260 11260 11260 11260 11260 ...
## $ Totalnsites  : int  2439 2439 2439 2439 2439 2439 2439 2439 2439 2439 ...
## $ TotalvisitPerYear: num  240 240 240 240 240 ...
## $ nyears       : int  47 47 47 47 47 47 47 47 47 ...
## $ FirstYrConverged : logi  FALSE TRUE TRUE TRUE TRUE ...
## $ LastYrConverged : logi  FALSE TRUE TRUE TRUE TRUE ...
## $ PropYrConverged : num  0 1 1 1 1 1 0 1 1 ...

head(model_data)

##           speciesName Spnsite Spnvisits SpvisitPerSite SpvisitPerYear
## 1 FORMICA aquilonia     51      94    1.843137  2.0000000
## 2 FORMICA cunicularia   336     785    2.336310 16.7021277
## 3 FORMICA exsecta       16      35    2.187500  0.7446809
## 4 FORMICA fusca        852    1804    2.117371 38.3829787
## 5 FORMICA lemani       237     450    1.898734  9.5744681
## 6 FORMICA lugubris      75     168    2.240000  3.5744681

```

```

##   Totalnvisits Totalnsites TotalvisitPerYear nyears FirstYrConverged
## 1      11260      2439      239.5745     47      FALSE
## 2      11260      2439      239.5745     47      TRUE
## 3      11260      2439      239.5745     47      TRUE
## 4      11260      2439      239.5745     47      TRUE
## 5      11260      2439      239.5745     47      TRUE
## 6      11260      2439      239.5745     47      TRUE
##   LastYrConverged PropYrConverged
## 1            FALSE              0
## 2             TRUE              1
## 3             TRUE              1
## 4             TRUE              1
## 5             TRUE              1
## 6             TRUE              1

```

Posteriors

We can also get data on the trend estimate from the 1000 posteriors that Charlie has already calculated. This is for all species groups regardless of whether they were run on Cirrus or JASMIN. We have to remove some years full of NAs. This is where the model is run to a year where there is no data. As a consequence some of the ‘final 10 years’ runs will actually be for a shorter period.

```

path_to_posteriors <-
  file.path('W:/PYWELL_SHARED/Pywell Projects/BRC/Charlie/1.c. New Model Rerun',
            '6. Indicators and other analyses/1000 posterior samples')

rdata_files <- list.files(path = tolower(path_to_posteriors),
                           pattern = '.rdata$',
                           full.names = TRUE)

annual_growth_rate <- function(years){
  years <- na.omit(years)
  (((tail(years, 1)/years[1])^(1/length(years))-1)*100
}

modelposterior <- function(sp, mat_j_post, sp_list, suffix = NULL){

  cat('Modelling', sp, '...\n')

  gr1000 <- apply(X = mat_j_post[sp_list == sp, ],
                  MARGIN = 1,
                  FUN = annual_growth_rate)

  mean_growth_rate <- mean(gr1000)
  precision_growth_rate <- 1/(sd(gr1000)^2)

  yrprec <- apply(X = mat_j_post[sp_list == sp, ],
                  MARGIN = 2,
                  FUN = function(x){
                    1/(sd(x)^2)
                  })

  mean_year_precision <- mean(yrprec, na.rm = TRUE)
  return(data.frame(species = paste0(sp, suffix),

```

```

        mean_growth_rate,
        precision_growth_rate,
        mean_year_precision))
}

for(data_file in rdata_files){

  # data_file <- rdata_files[1]
  cat('Starting', basename(data_file))

  # Charlie has named some of the objects differently, so I deal with this
  if(exists('j_post')) rm(list = 'j_post')
  if(exists('samp_post')) rm(list = 'samp_post')

  load(data_file)
  if(exists('j_post')){
    samp_post <- j_post
    rm(list = 'j_post')
  }

  # sp_list <- j_post$spp
  sp_list <- samp_post$spp

  spp <- unique(sp_list)

  cat('\n', length(spp), 'species\n')

  # mat_j_post <- as.matrix(j_post[, colnames(j_post)[!colnames(j_post) %in% c('iter', 'spp')]])
  mat_j_post <-
    as.matrix(samp_post[, colnames(samp_post)[!colnames(samp_post) %in% c('iter', 'spp')]])

  ## NOTE ##
  # Some of the posterior files contain data for years that the model was not run for
  # these appear here as columns of entirely NA values. These should be removed
  good_columns <- apply(mat_j_post, MARGIN = 2, FUN = function(x) !all(is.na(x)))
  if(any(!good_columns)){
    warning('Removing NA year(s) in ', basename(data_file), ' : ',
           paste(names(good_columns[!good_columns]), collapse = ', '))
    mat_j_post <- mat_j_post[,good_columns]
  }

  # set year estimates of 0 to 0.0001 to avoid infinite growth rates
  mat_j_post[mat_j_post == 0] <- 0.0001
  # rm(list = 'j_post')
  rm(list = 'samp_post')

  df <- lapply(spp, FUN = modelposterior, mat_j_post = mat_j_post, sp_list = sp_list)
  df <- do.call(rbind, df)
  row.names(df) <- 1:nrow(df)
  write.csv(df,
            row.names = FALSE,
            file = file.path('Results/metrics',
                             paste0('posteriorLM_',

```

```

        gsub('.rdata$', '',
              basename(data_file)),
              '.csv')))

df <- lapply(spp, FUN = modelposterior,
             mat_j_post = mat_j_post[, (ncol(mat_j_post)-9):ncol(mat_j_post)],
             sp_list = sp_list, suffix = '_last10yrs')
df <- do.call(rbind, df)
row.names(df) <- 1:nrow(df)
write.csv(df, file = file.path('Results/metrics',
                               paste0('posteriorLM_last10yrs_',
                                      gsub('.rdata$', '',
                                            basename(data_file)),
                                      '.csv'))))
}

}

```

Once we have the data for all species we can add the data together into one file

```

LMfiles <- list.files(path = 'Results/metrics',
                      pattern = '^posteriorLM',
                      full.names = TRUE)

master <- NULL

for(i in LMfiles){

  x <- read.csv(i)
  master <- rbind(master, x[,c('species',
                                'mean_growth_rate',
                                'precision_growth_rate',
                                'mean_year_precision')])

}

write.csv(master, file = 'Results/metrics/ALL_posteriorLM.csv',
          row.names = FALSE)

```

Raw data files

There are some additional stats we can get from the raw data file. Again, these are for all groups not just CIRRUS/JASMIN runs. Given the memory requirements needed to reformat the plant and moth data I have removed them from this. We can work with them separately by using the visit data files which are on JASMIN where they will have been created.

```

dataMetrics <- function(sp, basen, species_obs, hab=NULL,
                       habitat=FALSE, region=FALSE, suffix=NULL, years = FALSE){

  cat(as.character(sp), '... ')

  timetaken <- system.time({
    if(habitat){
      tFDall <- species_obs[(species_obs[,paste0(hab, '_hab')]),]
    } else if(region){
      tFDall <- species_obs[species_obs$nutsname==hab,]
    }
  })
}

```

```

} else {
  tFDall <- species_obs
}
tFD <- tFDall[tFDall[,names(tFDall) == sp],]

# Extract taxonomic name
tax_root <- substr(basen,start = 1,stop = (regexpr('_17',basen)[1]-1))

if(!is.null(years)){
  # There is a limit to the number of year's data we want
  # Set taxonomic name with this information included
  tax_nom <- paste0('last',years,'yrs',tax_root)
  # Drop everything we don't want
  if(nrow(tFD)!=0){
    boundingyear <- max(tFD$year)-years
    tFD <- tFD[tFD$year>boundingyear,]
  }
  else
  {
    # There's no observances for this species, so set a bounding year
    # based on the group data
    boundingyear <- max(tFDall$year)-years
  }
  tFDall <- tFDall[tFDall$year>boundingyear,]
} else {
  tax_nom <- tax_root
}

if(nrow(tFD)==0){
# No data, so set a few variables to 0. The reason for using 0's as defaults is if
# there is no data, these parameters should be 0 e.g. median number of visits = 0.
# Without this default, many parameters return NA, which is less accurate.
  nyears <- Spnsite <- Spnvisits <- SpvisitPerSite <- SpvisitPerYear <-
    avVisitPerYear <- median <- P70 <- P80 <- P90 <- sdVisitsPerYear <-
    coeffVar <- zmedian <- zP70 <- zP80 <- zP90 <- zsdVisitsPerYear <-
    zcoeffVar <- visits_median <- visits_P70 <- visits_P80 <- visits_P90 <-
    prop_repeats_spc <- prop_repeats_grp <- prop_of_years <- prop_list_one <-
    prop_abs_list <- Totalnvisits <- Totalnsites <- TotalvisitPerYear <- 0
  prop_abs <- 1
} else {
  # range of years with data
  nyears <- (1+max(tFD$year))-min(tFD$year))
  # Number of sites species is observed at
  Spnsite <- length(unique(tFD$site))
  # Number of visits on which the species is observed
  Spnvisits <- length(unique(tFD$visit))
  # Number of visits per site
  SpvisitPerSite <- Spnvisits/Spnsite
  # Number of visits per year
  SpvisitPerYear <- Spnvisits/nyears

  # average visits per year
  visits_year <- tapply(tFD$visit, tFD$year,

```

```

        FUN = function(x) length(unique(x)))
avVisitPerYear = mean(visits_year)

# find 70, 80 and 90th percentiles
percentiles <- quantile(visits_year,c(.5,.7,.8,.9))
median <- as.numeric(percentiles[1])
P70 <- as.numeric(percentiles[2])
P80 <- as.numeric(percentiles[3])
P90 <- as.numeric(percentiles[4])

if(nrow(tFD)==1){
  sdVisitsPerYear <- coeffVar <- 1
} else {
  # sd visits per year
  sdVisitsPerYear = sd(visits_year)
  # coefficient of variation
  coeffVar <- sd(visits_year)/mean(visits_year)
}

unique_years <- unique(tFD$year)
all_years     <- min(tFDall$year):max(tFDall$year)
missing_years <- all_years[!(all_years %in% unique_years)]
visits_year <- c(visits_year,rep(0,length(missing_years)))

# find 70, 80 and 90th percentiles
percentiles <- quantile(visits_year,c(.5,.7,.8,.9))
zmedian <- as.numeric(percentiles[1])
zP70 <- as.numeric(percentiles[2])
zP80 <- as.numeric(percentiles[3])
zP90 <- as.numeric(percentiles[4])

# sd visits per year
zsdVisitsPerYear = sd(visits_year)
# coefficient of variation
zcoeffVar <- sd(visits_year)/mean(visits_year)

# repeat visits
# Within each year get the counts of visits to each location
repeats <- count(tFD, site, year)
repeats$concat <- paste0(repeats$site,'-',repeats$year)

# What proportion of these are > 1
prop_repeats_spc <- sum(repeats$n > 1) / nrow(repeats)

# visits within the group for visits to each location within a year
group_repeats <- count(tFDall, site, year)
group_repeats$concat <-
  paste0(group_repeats$site,'-',group_repeats$year)

# Find which of these group visits are to sites in years where the
# species of interest was observed
site_year <- group_repeats$concat %in% repeats$concat
group_repeats <- group_repeats$n[site_year]

```

```

prop_repeats_grp <- sum(group_repeats > 1) / length(group_repeats)

# all visits to a site in a year where there was at least one observance
# of species of interest
percentiles <- quantile(group_repeats,c(.5,.7,.8,.9))
visits_median <- as.numeric(percentiles[1])
visits_P70 <- as.numeric(percentiles[2])
visits_P80 <- as.numeric(percentiles[3])
visits_P90 <- as.numeric(percentiles[4])

# proportion of years with data
prop_of_years <-
  length(unique(tFD$year))/(1+max(tFDall$year)-min(tFDall$year))

# lists of length 1
prop_list_one <- sum(tFD$L == 1) / nrow(tFD)
}

# proportion of records for the group which did not observe this species
prop_abs <- (nrow(tFDall)-nrow(tFD))/(nrow(tFDall))

# proportion of non-observances with list length > 1
tFDabs <- tFDall[!tFDall[,colnames(tFDall) %in% sp],]
prop_abs_list <- sum(tFDabs$L > 1) / nrow(tFDabs)

# Number of sites for taxonomic group
Totalnsites <- length(unique(tFDall$site))
# Number of visits for taxonomic group
Totalnvisits <- nrow(tFDall)
if(nrow(tFDall)==0){
  total_years <- 0
  TotalvisitPerYear <- 0
} else {
  total_years <- length(min(tFDall$year):max(tFDall$year))
  TotalvisitPerYear <- Totalnvisits/total_years
}

df <- data.frame(species = paste0(sp, suffix),
                  habitat = NA,
                  region = NA,
                  avVisitPerYear = avVisitPerYear,
                  median = median,
                  P70 = P70,
                  P80 = P80,
                  P90 = P90,
                  sdVisitsPerYear = sdVisitsPerYear,
                  coeffVar = coeffVar,
                  zmedian = zmedian,
                  zP70 = zP70,
                  zP80 = zP80,
                  zP90 = zP90,
                  zsdVisitsPerYear = zsdVisitsPerYear,
                  zcoeffVar = zcoeffVar,

```

```

prop_repeats_spc = prop_repeats_spc,
prop_repeats_grp = prop_repeats_grp,
visits_median = visits_median,
visits_P70 = visits_P70,
visits_P80 = visits_P80,
visits_P90 = visits_P90,
prop_of_years = prop_of_years,
prop_list_one = prop_list_one,
prop_abs = prop_abs,
prop_abs_list = prop_abs_list,
nyears = nyyears,
Spnsite = Spnsite,
Spnvisits = Spnvisits,
SpvisitPerSite = SpvisitPerSite,
SpvisitPerYear = SpvisitPerYear,
Totalnsites = Totalnsites,
Totalnvisits = Totalnvisits,
TotalvisitPerYear = TotalvisitPerYear,
Taxa_Root = tax_root,
Taxa = tax_nom,
stringsAsFactors = FALSE)

if(habitat){
  df[, 'habitat'] <- hab
} else if(region){
  df[, 'region'] <- hab
}
})
cat(as.numeric(timetaken[3]), 'seconds', '\n')
return(df)
}

data_files_path <- file.path('W:/PYWELL_SHARED/Pywell Projects/BRC/Charlie',
                             '1.c. New Model Rerun/1. Data/Cleaned Datasets')

data_files <- list.files(data_files_path,
                        pattern = '.rdata$',
                        full.names = TRUE)

# Dont do the really big ones!
data_files <- data_files[!grep('Moths', data_files)]
data_files <- data_files[!grep('VascPlants', data_files)]

dataPrep <- function(data_file){
  cat('Starting', basename(data_file))

  load(data_file)

  if('SQ_1KM' %in% colnames(taxa_data)){
    names(taxa_data)[names(taxa_data) == 'SQ_1KM'] <- 'TO_GRIDREF'
  }

  # Filter the data as it is for the occupancy models
  formatted_data <- formatOccData(taxa = taxa_data$CONCEPT,

```

```

        site = taxa_data$TO_GRIDREF,
        time_period = taxa_data$TO_STARTDATE)
    return(formatted_data)
}

for(data_file in data_files){
  formatted_data <- dataPrep(data_file)

  spp <- unique(taxa_data$CONCEPT)
  # one group has a species called '', best to remove this (single) record
  spp <- spp[spp != '']

  cat('\n', length(spp), 'species\n')

  tFDall <- merge(formatted_data$occDetdata,
                  formatted_data$spp_vis)

  # Remove sites which have not seen a repeat in subsequent years.
  # This is a requirement here as this was a step taken by the model.
  yps <- rowSums(acast(tFDall, site ~ year, length, value.var = 'L') > 0)
  sites_to_include <- names(yps[yps >= 2])
  tFDall <- tFDall[as.character(tFDall$site) %in% sites_to_include,]

  raw_metrics <- lapply(spp,
                        FUN = dataMetrics,
                        species_obs = tFDall,
                        basen = basename(data_file))

  raw_metrics <- do.call(rbind, raw_metrics)

  write.csv(raw_metrics, file = file.path('Results/metrics',
                                         paste0('rawMetrics_',
                                                gsub('.rdata$', '',
                                                      basename(data_file)),
                                                '.csv')),
            row.names = FALSE)

  # now with just the last 10 years
  cat('\n', length(spp), 'species\n')
  raw_metrics <- lapply(spp,
                        FUN = dataMetrics,
                        tFDall = tFDall,
                        basen = basename(data_file),
                        suffix = '_last10yrs',
                        years = 10)

  raw_metrics <- do.call(rbind, raw_metrics)

  write.csv(raw_metrics, file = file.path('Results/metrics',
                                         paste0('rawMetrics_last10yrs',
                                                gsub('.rdata$', '',
                                                      basename(data_file)),
                                                '.csv'))),

```

```

    row.names = FALSE)
}

```

Once we have the data for all species we can add the data together into one file

```

rawdataFiles <- list.files(path = 'Results/metrics',
                           pattern = '^rawMetrics',
                           full.names = TRUE)

master <- NULL

for(i in rawdataFiles){

  x <- read.csv(i)
  master <- rbind(master, x)

}

write.csv(master, file = 'Results/metrics/ALL_rawMetrics.csv',
          row.names = FALSE)

```

Classification of Models

Create metrics

We can do some simple decision tree statistics to develop rules of thumb once we have decided what constitutes a ‘bad’ model. Here I choose some metrics that make for a bad model and use the `rpart` package to create a decision tree.

First, combine all the data

```

RM <- read.csv('Results/metrics/ALL_rawMetrics.csv')
LM <- read.csv('Results/metrics/ALL_posteriorLM.csv')
MM <- read.csv('Results/metrics/model_data.csv')

# Remove repeat variables from model data
MM <- MM[,c('speciesName','FirstYrConverged',
           'LastYrConverged','PropYrConverged')]

# These merges drop a LOT of species that don't match. That is OK as the ones that don't
# match do not meet the data minimum criteria or we do not have model data for them, so
# cannot use them for predictions.
trendsData <- merge(x = LM, y = RM,
                     by.x = 'species',
                     by.y = 'species')
trendsData <- merge(x = MM, y = trendsData,
                     by.x= 'speciesName', by.y = 'species')
head(trendsData)

##                                     speciesName FirstYrConverged LastYrConverged
## 1             Lamprochromus bifasciatus      TRUE            TRUE
## 2 Lamprochromus bifasciatus_last10yrs      TRUE            TRUE
## 3           Acanthosoma haemorrhoidale      TRUE            TRUE
## 4           Achalcus bimaculatus      TRUE            TRUE
## 5 Achalcus bimaculatus_last10yrs     FALSE            TRUE

```

```

## 6          Achalcus britannicus      TRUE      FALSE
## PropYrConverged mean_growth_rate precision_growth_rate
## 1      1.0000000    2.9382204    0.024610675
## 2      1.0000000   -12.0366451    0.002793217
## 3      1.0000000    0.1742559    0.847053523
## 4      0.6382979    0.2458987    0.610197450
## 5      0.8000000    0.5830347    0.055452645
## 6      0.7234043    1.3467195    0.047072983
## mean_year_precision avVisitPerYear Spnvisits Spnsite SpvisitPerSite
## 1      999.99334    3.571429    25     18    1.388889
## 2      83.24288    5.000000    20     14    1.428571
## 3      158.84711   26.534884   1141    760    1.501316
## 4      21.48738    1.000000     3      3    1.000000
## 5      21.36882    1.000000     1      1    1.000000
## 6      86108.00740  7.833333    47      5    9.400000
## SpvisitPerYear Totalnvisits Totalnsites TotalvisitPerYear median P70
## 1      0.8928571    13201     2135    280.8723   2.0  3.8
## 2      3.3333333    4765      1172    297.8125   5.0  7.2
## 3      24.8043478   10914     2441    232.2128  12.0 36.8
## 4      0.1153846    13201     2135    280.8723   1.0  1.0
## 5      1.0000000    3650      949    280.7692   1.0  1.0
## 6      3.9166667    13201     2135    280.8723   4.5  9.5
## P80 P90 sdVisitsPerYear coeffVar zmedian zP70 zP80 zP90
## 1  6.2 7.8    3.154739  0.8833270    0    0  0.0  1.4
## 2  7.8 8.4    3.651484  0.7302967    0    0  1.0  5.0
## 3 52.8 69.0   29.763631  1.1216794   11   28 50.8 69.0
## 4  1.0 1.0    0.000000  0.0000000    0    0  0.0  0.0
## 5  1.0 1.0     NA       NA      0    0  0.0  0.0
## 6 13.0 18.0   8.681398  1.1082635    0    0  0.0  1.0
## zsdVisitsPerYear zcoeffVar prop_repeats_spc prop_repeats_grp
## 1      1.7175699  3.229031    0.20000000  0.6000000
## 2      2.7688746  2.215100    0.18750000  0.5625000
## 3      29.4084446  1.211391    0.08442211  0.3939698
## 4      0.2470922  3.871112    0.00000000  1.0000000
## 5      0.2773501  3.605551    0.00000000  1.0000000
## 6      3.8953707  3.895371    0.71428571  0.8571429
## visits_median visits_P70 visits_P80 visits_P90 prop_of_years
## 1          2      3.0      3.4      6.0  0.14893617
## 2          2      3.0      3.0      6.0  0.25000000
## 3          1      2.0      3.0      5.0  0.91489362
## 4          2      9.6     13.4     17.2  0.06382979
## 5         21     21.0     21.0     21.0  0.07692308
## 6          9     21.8     24.2     27.8  0.12765957
## prop_abs_list prop_abs prop_list_one      Taxa Taxa_Root
## 1      0.5651943  0.9981062    0.08000000    E&D    E&D
## 2      0.6307692  0.9958027    0.00000000 last10yrsE&D    E&D
## 3      0.2566254  0.8954554    0.52673094 ShieldBugs ShieldBugs
## 4      0.5658433  0.9997727    0.33333333    E&D    E&D
## 5      0.6440121  0.9997260    1.00000000 last10yrsE&D    E&D
## 6      0.5645431  0.9964397    0.06382979    E&D    E&D

str(trendsData)

## 'data.frame': 9323 obs. of 39 variables:
## $ speciesName : Factor w/ 9266 levels "", "Lamprochromus bifasciatus", ... : 2 3 4 5 6 7 8 9 ...

```

```

## $ FirstYrConverged      : logi  TRUE TRUE TRUE TRUE FALSE TRUE ...
## $ LastYrConverged       : logi  TRUE TRUE TRUE TRUE TRUE FALSE ...
## $ PropYrConverged        : num   1 1 1 0.638 0.8 ...
## $ mean_growth_rate       : num   2.938 -12.037 0.174 0.246 0.583 ...
## $ precision_growth_rate: num   0.02461 0.00279 0.84705 0.6102 0.05545 ...
## $ mean_year_precision   : num   1000 83.2 158.8 21.5 21.4 ...
## $ avVisitPerYear         : num   3.57 5 26.53 1 1 ...
## $ Spnvisits              : int   25 20 1141 3 1 47 11 128 71 62 ...
## $ Spnsite                : int   18 14 760 3 1 5 4 66 34 45 ...
## $ SpvisitPerSite          : num   1.39 1.43 1.5 1 1 ...
## $ SpvisitPerYear          : num   0.893 3.333 24.804 0.115 1 ...
## $ Totalnvisits            : int   13201 4765 10914 13201 3650 13201 7115 13201 3463 13201 ...
## $ Totalnsites              : int   2135 1172 2441 2135 949 2135 1602 2135 888 2135 ...
## $ TotalvisitPerYear        : num   281 298 232 281 281 ...
## $ median                  : num   2 5 12 1 1 4.5 2 2 2 2 ...
## $ P70                      : num   3.8 7.2 36.8 1 1 9.5 3.3 3.6 5.6 3.4 ...
## $ P80                      : num   6.2 7.8 52.8 1 1 13 4.2 5 14.6 4 ...
## $ P90                      : num   7.8 8.4 69 1 1 18 5.1 10.8 18.2 5 ...
## $ sdVisitsPerYear          : num   3.15 3.65 29.76 0 NA ...
## $ coeffVar                 : num   0.883 0.73 1.122 0 NA ...
## $ zmedian                  : num   0 0 11 0 0 0 0 1 2 0 ...
## $ zP70                     : num   0 0 28 0 0 0 0 2 2 1 ...
## $ zP80                     : num   0 1 50.8 0 0 0 0 3 11.6 2 ...
## $ zP90                     : num   1.4 5 69 0 0 1 1 5.4 16.7 4 ...
## $ zsdVisitsPerYear          : num   1.718 2.769 29.408 0.247 0.277 ...
## $ zcoeffVar                : num   3.23 2.22 1.21 3.87 3.61 ...
## $ prop_repeats_spc          : num   0.2 0.1875 0.0844 0 0 ...
## $ prop_repeats_grp          : num   0.6 0.562 0.394 1 1 ...
## $ visits_median             : num   2 2 1 2 21 9 6.5 2 2 2 ...
## $ visits_P70                : num   3 3 2 9.6 21 21.8 8.1 2 2 3 ...
## $ visits_P80                : num   3.4 3 3 13.4 21 ...
## $ visits_P90                : num   6 6 5 17.2 21 ...
## $ prop_of_years              : num   0.1489 0.25 0.9149 0.0638 0.0769 ...
## $ prop_abs_list              : num   0.565 0.631 0.257 0.566 0.644 ...
## $ prop_abs                  : num   0.998 0.996 0.895 1 1 ...
## $ prop_list_one              : num   0.08 0 0.527 0.333 1 ...
## $ Taxa                      : Factor w/ 64 levels "Ants","AquaticBugs",...: 9 25 58 9 25 9 25 9 25 9 ...
## $ Taxa_Root                 : Factor w/ 32 levels "Ants","AquaticBugs",...: 9 9 26 9 9 9 9 9 9 9 ...

```

Consultation on model outputs

Next, define which of the species have ‘good’ models and which are ‘bad’.

This was done by consultation with 3 experts on the models, testing them on 100 models which were sampled to provide a spread of examples across the dataset, with a cluster of models focused around the range in which there was likely to be controversial.

```
# Lets have a look at the distribution  
quantile(trendsData$precision_growth_rate)
```

```
##          0%        25%        50%        75%       100%  
## 0.000582423 0.054378193 0.153248562 0.472834262 18.948317680
```

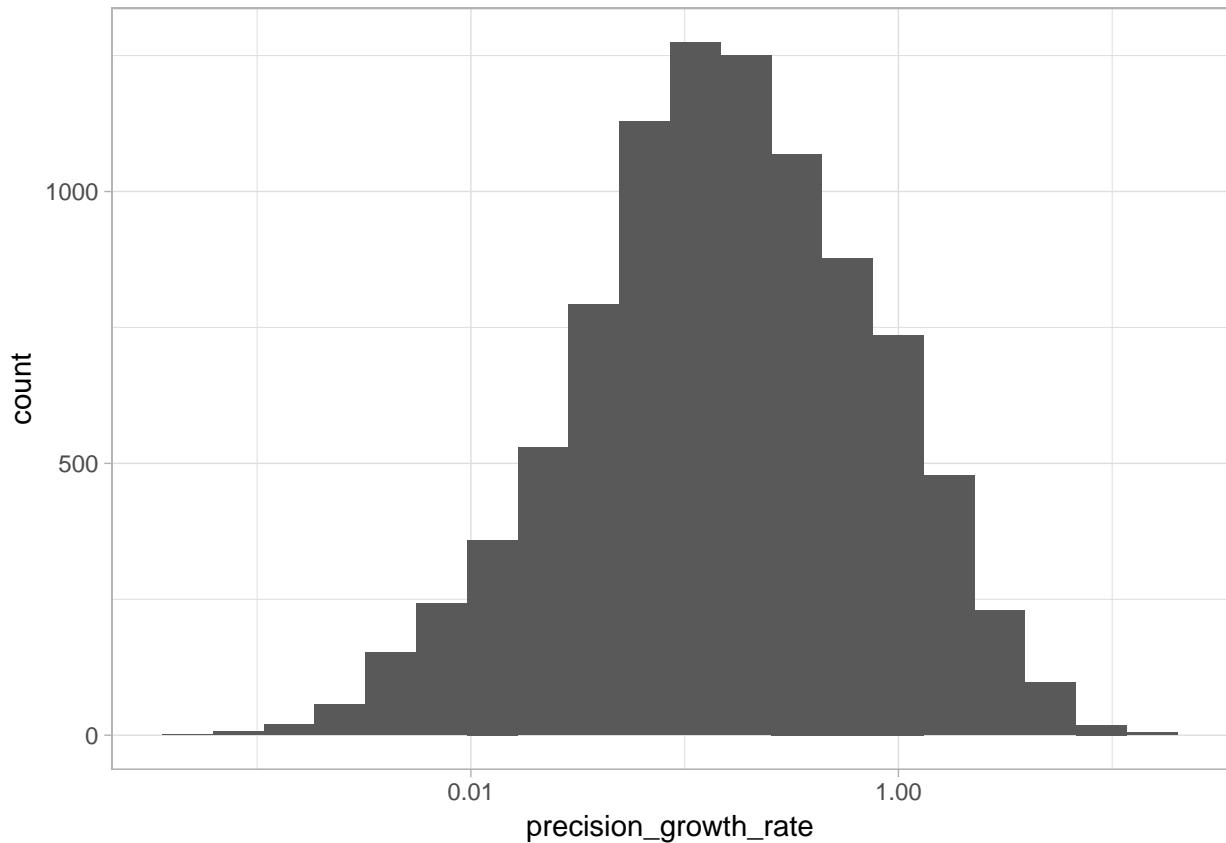


Figure 1: Histogram of the growth rate precision for all models

```
quantile(trendsData$mean_year_precision)
##          0%        25%        50%        75%       100%
## 4.721129e+00 2.451261e+01 2.197646e+02 1.443673e+03 8.739698e+06
```

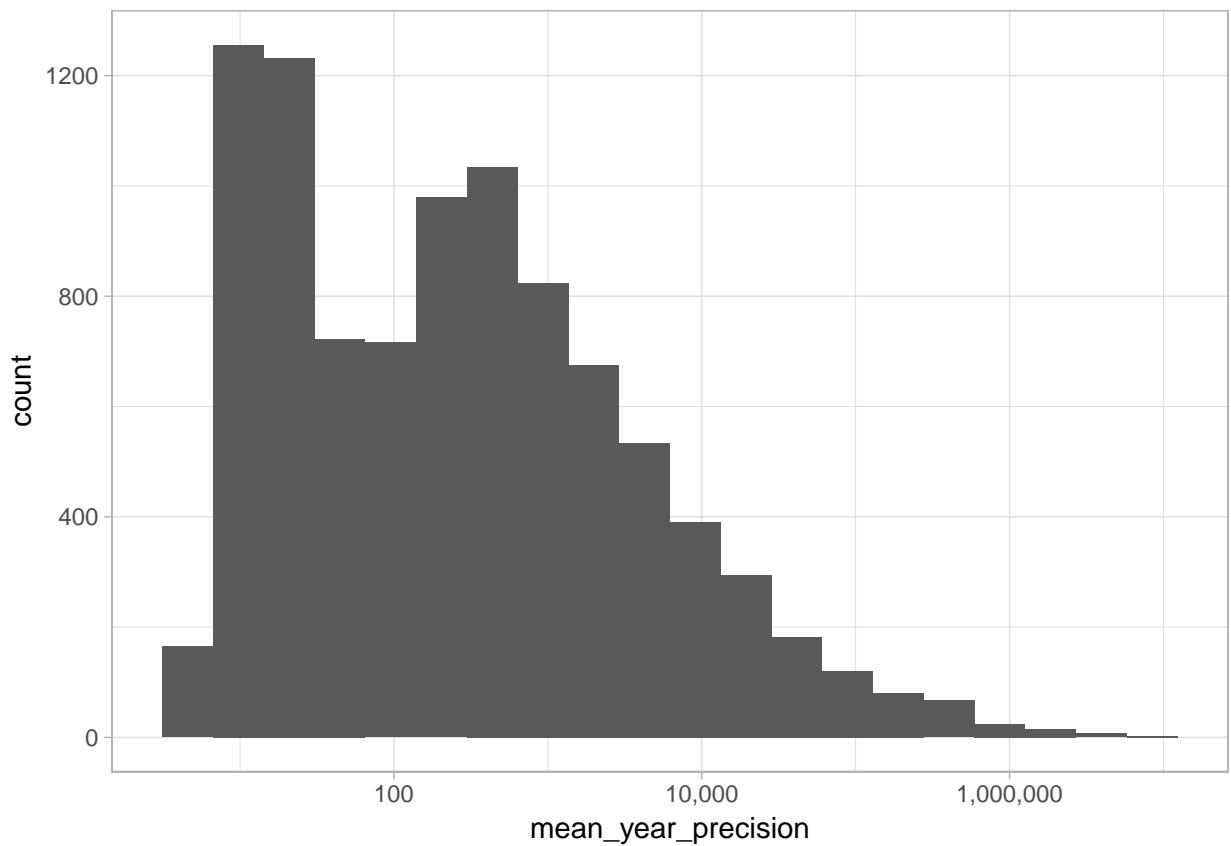


Figure 2: Histogram of the mean year precision for all models

Below are the results of the consultation. This graph shows the score given to each model where the score is the number of experts who thought the model was of good quality. The models not used in the assessment are shown in the background in grey

```
# Lets look at the results of the consultation with model experts  
fsdf <- read.csv(file = 'Results/Consultation/fsdf_full.csv')
```

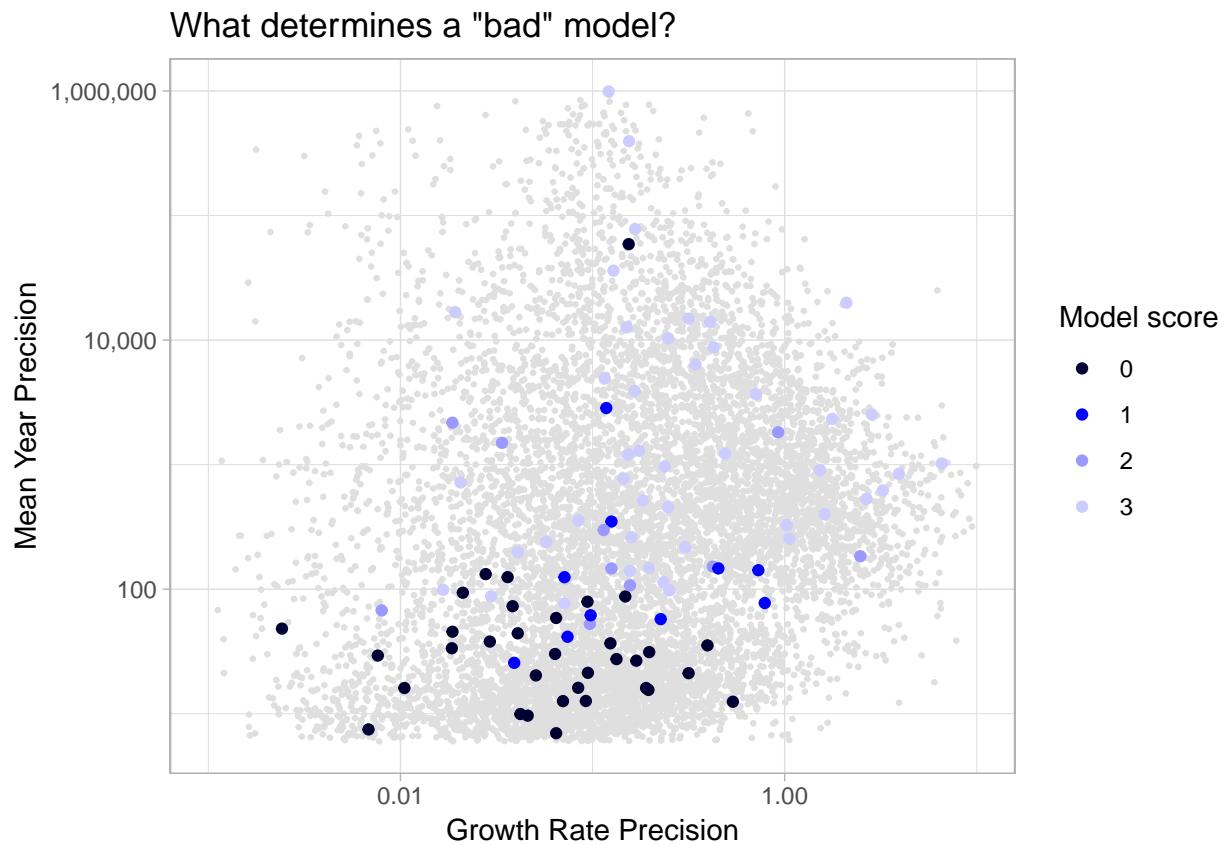


Figure 3: Graph showing results of consultation with model experts, plotting mean year precision vs growth rate precision

What determines a "bad" model?

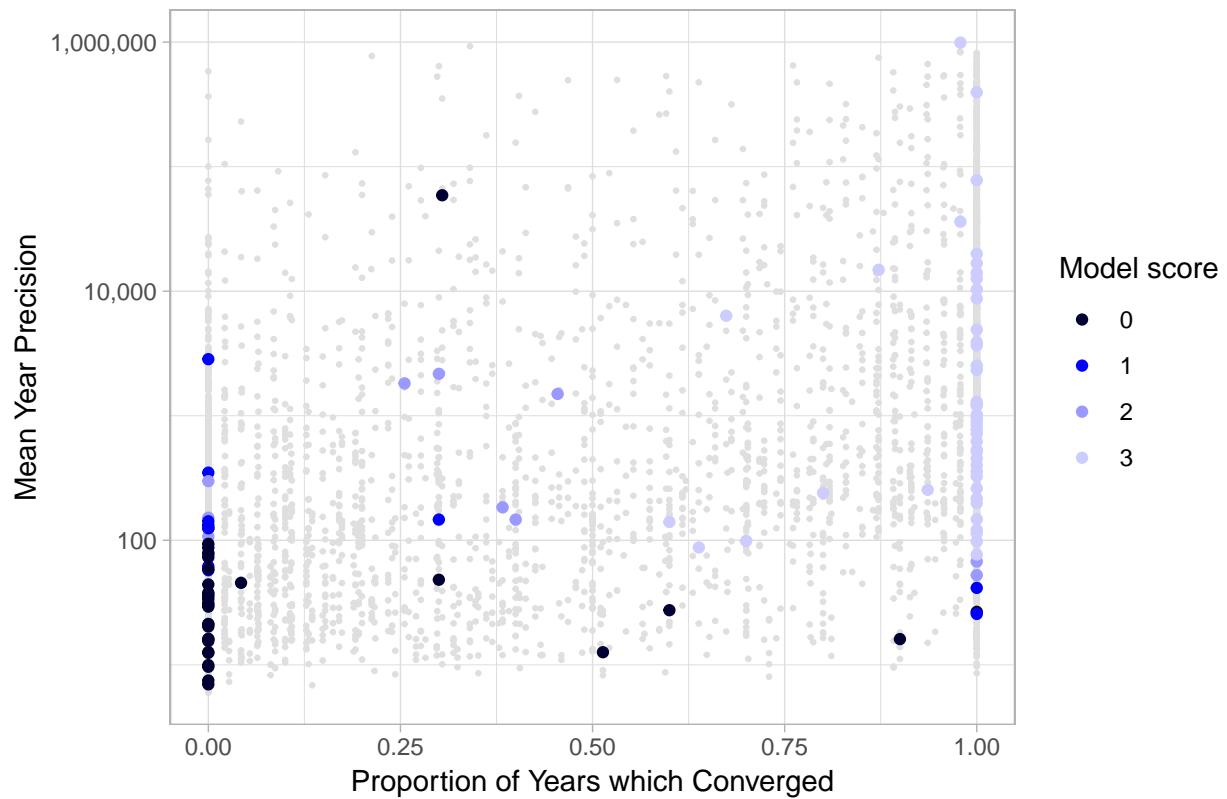


Figure 4: Graph showing results of consultation with model experts, plotting mean year precision vs the proportion of years which converged. Coloured points show the score assigned to models: light blue represents a score of 3 or unanimously a good model; black shows a score of 0, or unanimously a bad model. Light grey points in the background show all models which were not part of the consultation

The graphs above show the results of the consultation with experts on the model. There was unanimous agreement on 80% of the models about which were ‘good’ and ‘bad’. The models chosen for this consultation were deliberately difficult to tell apart, so a 80% unanimous agreement is considered very good.

Plotting these scores against proportion of years which converged, mean year precision and growth rate precision show that these parameters are correlated with quality, but it is not obvious how to perform this split.

Classification tree

To decide which models would be classified as good and bad, a decision tree was created to automatically classify the 100 models tested on. The models with a score of 2 or 3 were classified as good, while the models with a score of 0 or 1 were classified as bad.

Two decision trees are shown below, one using the proportion of years converged, and the other not, for comparison.

```
fsdf$good <- rep('bad', nrow(fsdf))
fsdf$good[fsdf$score>=2] <- 'good'

fit_fsdf <- rpart(good ~ mean_year_precision + precision_growth_rate,
                    method = 'class',
                    data = fsdf)

plot(fit_fsdf, uniform=TRUE,
      main = "Which models are good or bad (precision only)?",
      margin = .1)
text(fit_fsdf, use.n = TRUE,
     all = TRUE,
     cex = .7)

fit_fsdf <- rpart(good ~ mean_year_precision + precision_growth_rate + PropYrConverged,
                    method = 'class',
                    data = fsdf)

plot(fit_fsdf, uniform=TRUE,
      main = "Which models are good or bad (precision and convergence)?",
      margin = .1)
text(fit_fsdf, use.n = TRUE,
     all = TRUE,
     cex = .7)
```

Which models are good or bad (precision only)?

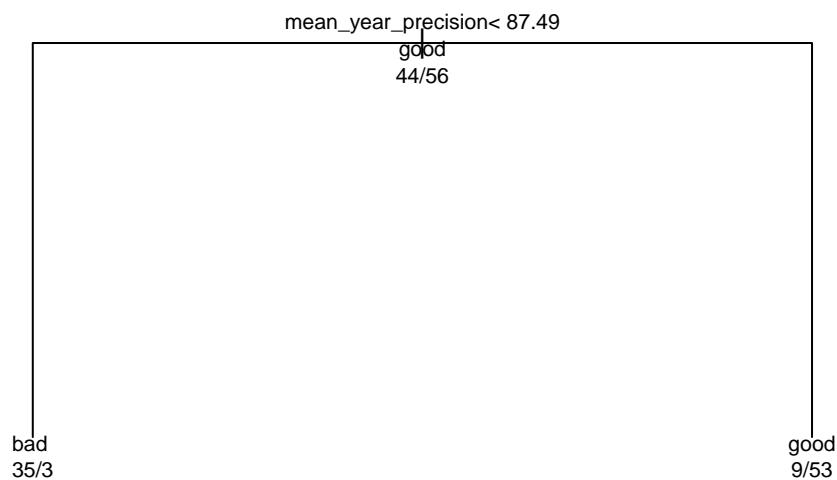


Figure 5: Decision tree for deciding which models are good or bad based on precision only

Which models are good or bad (precision and convergence)?

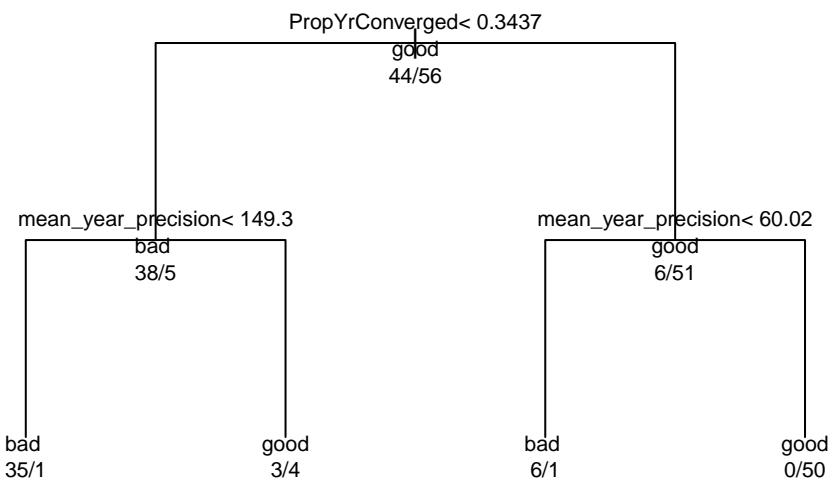


Figure 6: Decision tree for deciding which models are good or bad based on precision and convergence

The success of the decision tree for classifying bad models is:

- Using precision only: 92% for bad models, 85% for good models
- Using precision and convergences: 95% for bad models, 95% for good models

Using both metrics is clearly better for splitting good and bad models, therefore this split will be used. The split is as follows:

- If the proportion of years converged is < 0.3437 , the model will be classified as ‘bad’ if the mean year precision is < 149.3
- If the proportion of years converged is ≥ 0.3437 , the model will be classified as ‘bad’ if the mean year precision is < 60.2

Classifying the data

Examining the classification variables

Before running the classifiers, it is important to determine whether or not any of the variables are highly taxa dependent. If they are, they are much less generalisable across data sets.

This was done for all variables and the results from two of the variables are shown below:

- First is the violin plot for prop_repeats_grp: the proportion of site:year combinations for the species of interest which have more than one visit.
- Second is the violin plot for prop_abs_list: the proportion of data for the taxonomic group *not including the species of interest* which have a list length > 1 .

The prop_repeats_grp data shows a broad range of results from 0 to 1 for all taxonomic groups. By contrast, the prop_abs_list data shows very clear division between taxonomic groups. Therefore, any partitioning on the basis of this variable is effectively a taxonomic division. For this reason, prop_abs_list will not be used for producing decision trees.

```
# Subset data to remove last 10 yr data, as it just makes the plot messier
td_taxa <- trendsData[as.character(trendsData$Taxa)==
                        as.character(trendsData$Taxa_Root),]
```

Violin Plot of the proportion of repeats

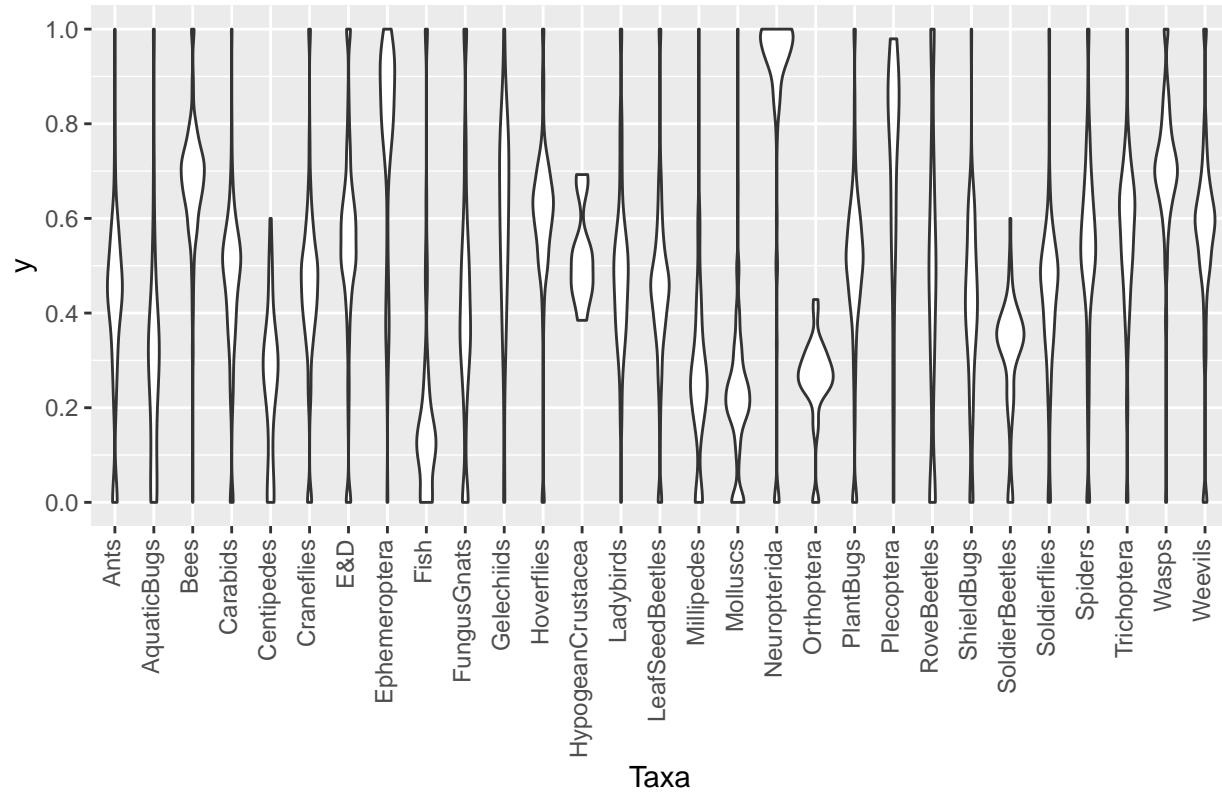


Figure 7: Violin plot for the prop_repeats_grp variable, defined as the proportion of all site;year combinations, where the focal species is observed, that have more than 1 visit for that taxonomic group

Violin Plot of proportion of absence data

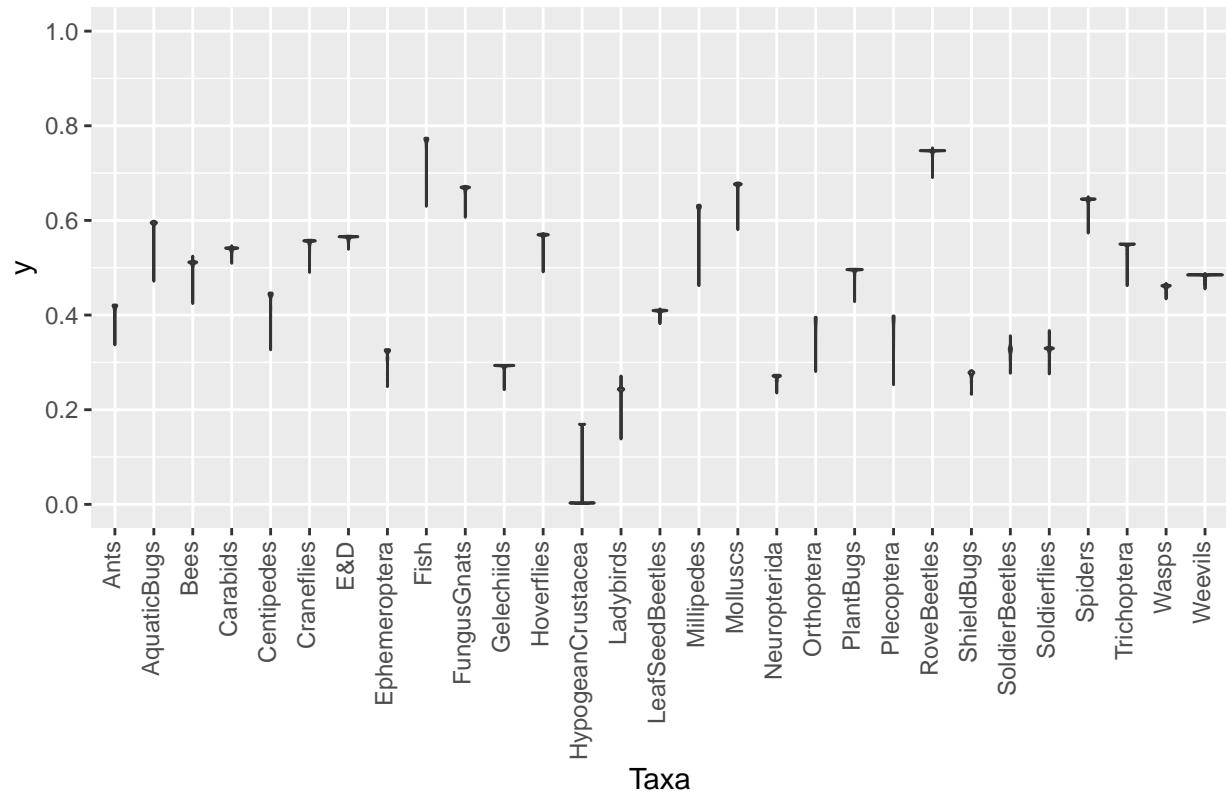


Figure 8: Violin plot for the prop_abs_list variable, defined as the proportion of records with list length greater than 1, for all records where the focal species was not observed

Correlation between variables

To examine the between all the variables, a pairs plot is produced below. As can be seen, P90 is strongly correlated with number of records (Spnvisits), as expected. Prop_abs is negatively correlated with number of records, as the fewer the records for a species, the higher the proportion of taxonomic data without the species of interest.

Apart from number of records, which is not included in the classification tree, the variables show little cross correlation, suggesting they are a good choice for building a decision tree.

```
td_sub <- trendsData[,c('P90','visits_P90','prop_of_years',
                      'prop_repeats_grp','prop_list_one',
                      'prop_abs','Spnvisits')]

pairs.panels(td_sub, hist.col = 'blue',smooth = TRUE)
```

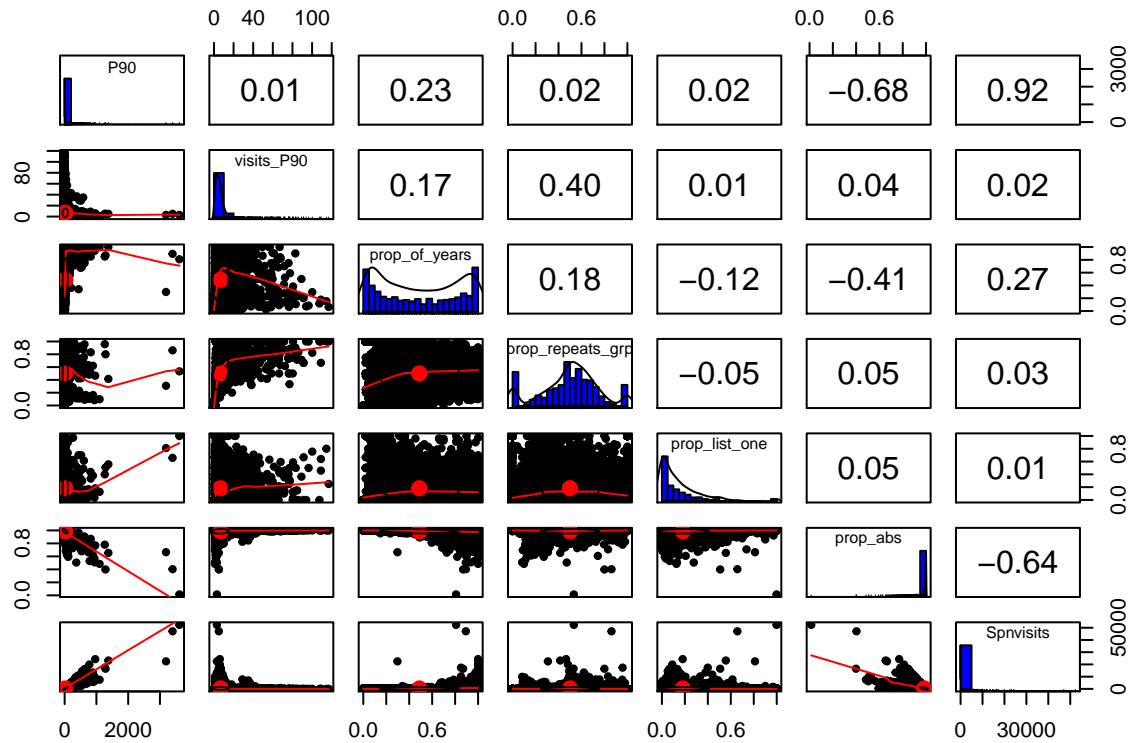


Figure 9: Pairs plot for all variables used to construct the classification tree, plus the Spnvisits variable, defined as the total number of records for each species of interest

Input data requirement

In order to be included in the dataset, one preliminary filtering step was carried out to remove some of the data. This step was to remove all records from visits to sites that never saw a repeat visit in any other years, for that taxonomic group. This step removed 33% of all records across our database.

- For instance, if a dragonfly was recorded in one site in 1990, if there was one more record of any dragonfly at that site in any year except 1990, both records would be included.
- If there was another record for another dragonfly to the same site in 1990, but no other records for that site, both of these records would be removed from the database.

This step was carried out as it was carried out on the data prior to running the models. Failing to do this step would make the results invalid.

This means that any decision tree below has the pre-requisite that all records in the database must meet this requirement.

To show the effect of this step, below is a plot showing the proportion of the records removed for each taxonomic group.

```
# Read in the metrics
RM <- read.csv('Results/metrics/ALL_rawMetrics.csv')
RM <- RM[RM$Taxa==as.character(RM$Taxa_Root),]
# Read in the metrics calculated for all excluded records
RM_1rec <- read.csv('Results/metrics/ALL_1rec.csv')
colnames(RM_1rec)[2] <- 'numrec_removed'
RM <- merge(RM,RM_1rec)
RM$Taxa <- as.character(RM$Taxa)
df <- NULL

# Calculate proportion of records removed for each taxonomic group
for(taxa in sort(unique(RM$Taxa))){
  num_inc <- sum(RM$Spnvisits[RM$Taxa==taxa])
  num_exc <- sum(RM$numrec_removed[RM$Taxa==taxa])
  df <- rbind(df,
    data.frame(taxa = taxa,
      Removed = (num_exc/(num_inc+num_exc)),
      Included = (num_inc/(num_inc+num_exc))))
}

taxa_melt <- melt(df, id=c('taxa'))

stack_records(taxa_melt, colours = c('red', '#9999FF'),
  ylabel = 'Proportion of records',
  title = 'Proportion of records which are included')
```

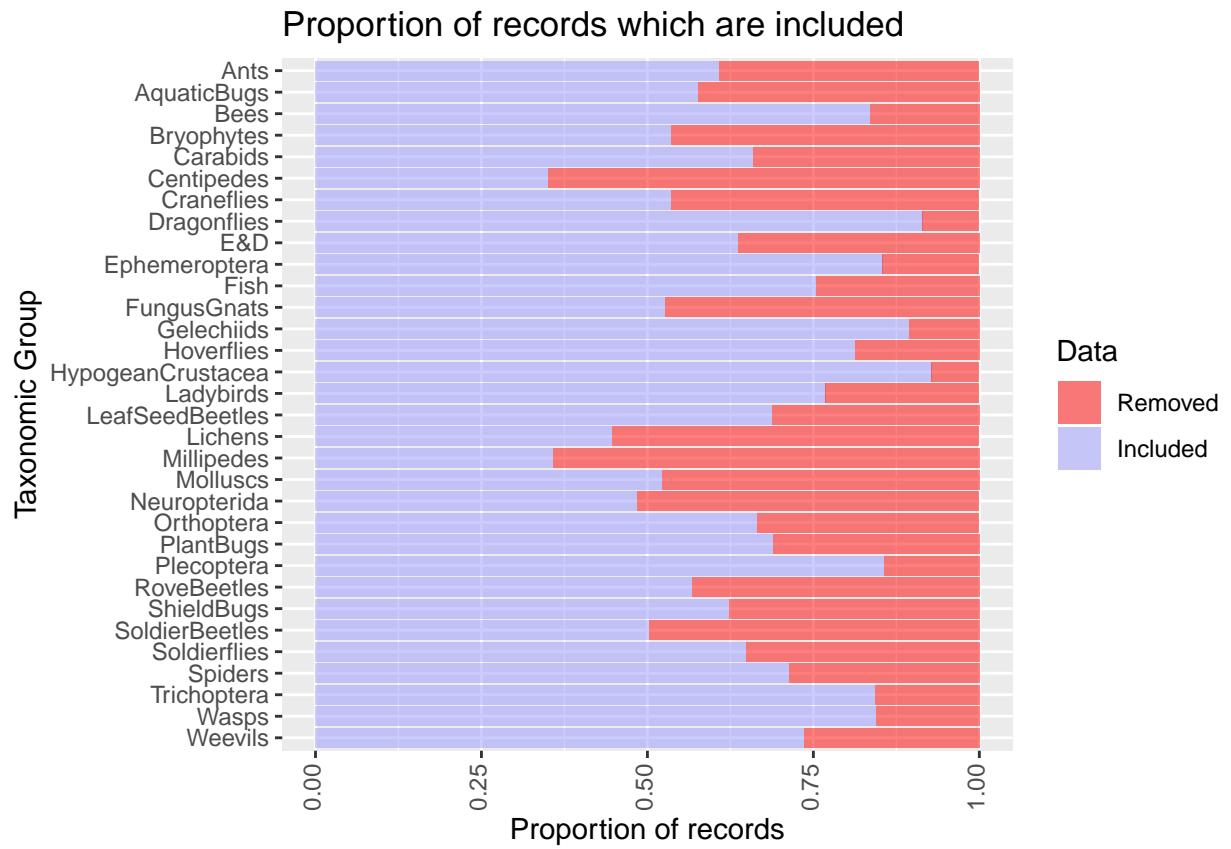


Figure 10: Graph showing the proportion of records which are included in the models. Those which are not are excluded or to sites which only received visit(s) in one year within the taxonomic group

Classification trees

Equally weighted decision tree

Using the split in the models between ‘good’ and ‘bad’, you can create a classification tree that attempts to use the variables we have extracted to partition the data to the best of its ability using a decision tree.

```
trendsData$bad <- rep('bad', nrow(trendsData))
trendsData$bad[(trendsData$PropYrConverged < 0.3437 &
               trendsData$mean_year_precision >= 149.3) | 
               (trendsData$PropYrConverged >= 0.3437 &
               trendsData$mean_year_precision >= 60.02)] <- 'good'

# Perform the fit. The loss parameter in the rpart function determines the
# weighting:
# By increasing 'goodweight', this increases the importance of correctly
# identifying good models and not assigning bad models to the good category.
# However, this is at the expense of assigning many good models to the bad
# category.
# By increasing 'badweight', the inverse is true
badweight <- 1
goodweight <- 1

fit <- rpart(bad ~ median + P90 + zmedian + zP90 +
              visits_median + visits_P70 + visits_P80 + visits_P90 +
              prop_of_years + prop_repeats_grp + prop_list_one +
              prop_abs,
              method = "class",
              data = trendsData,
              parms = list(loss = matrix(c(0,goodweight,badweight,0),ncol = 2)))

printcp(fit) # display the results

## 
## Classification tree:
## rpart(formula = bad ~ median + P90 + zmedian + zP90 + visits_median +
##       visits_P70 + visits_P80 + visits_P90 + prop_of_years + prop_repeats_grp +
##       prop_list_one + prop_abs, data = trendsData, method = "class",
##       parms = list(loss = matrix(c(0, goodweight, badweight, 0),
##                                   ncol = 2)))
##
## Variables actually used in tree construction:
## [1] P90      prop_abs
##
## Root node error: 3805/9323 = 0.40813
##
## n= 9323
##
##          CP nsplit rel error  xerror     xstd
## 1 0.577924      0    1.00000  1.00000 0.0124720
## 2 0.011038      1    0.42208  0.42234 0.0095845
## 3 0.010000      3    0.40000  0.41130 0.0094842
```

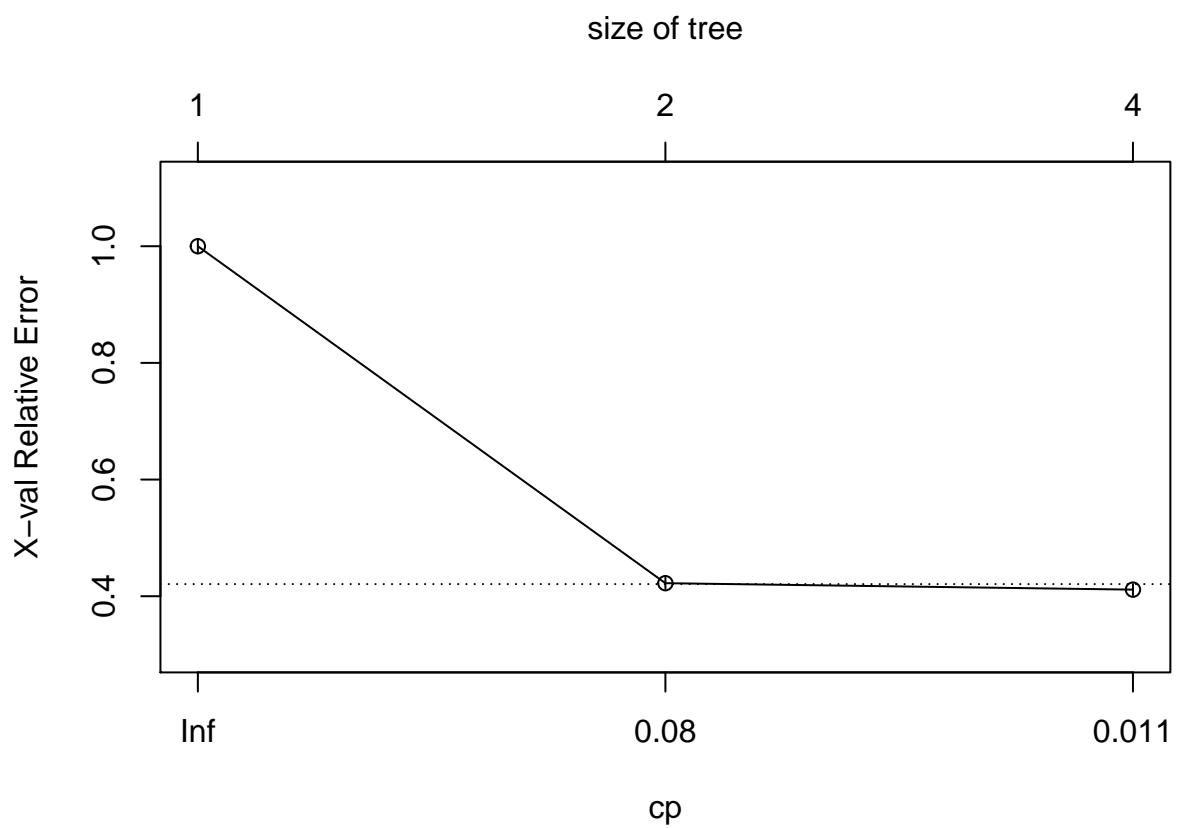


Figure 11: Visual representation of the results. Plotted is cross-validation error vs the complexity parameter

Below is a summary of the fit calculations. The relative importance of the variables are shown, which highlights that the most important variables are:

- P90: 90th percentile of records in a year for the species of interest
- median: median of records in a year for the species of interest
- prop_abs: proportion of visits within a taxonomic group which did not record the species of interest
- prop_of_years: proportion of years with records for the species of interest
- visits_P90: 90th percentile of records in a year, for sites where the species of interest was observed in that year
- prop_list_one: proportion of records for which the species of interest is recorded, with list length of 1

```
summary(fit) # detailed summary of splits
```

```
## Call:
## rpart(formula = bad ~ median + P90 + zmedian + zP90 + visits_median +
##        visits_P70 + visits_P80 + visits_P90 + prop_of_years + prop_repeats_grp +
##        prop_list_one + prop_abs, data = trendsData, method = "class",
##        parms = list(loss = matrix(c(0, goodweight, badweight, 0),
##                                    ncol = 2)))
## n= 9323
##
##          CP nsplit rel error      xerror       xstd
## 1 0.57792378      0 1.0000000 0.012471988
## 2 0.01103811      1 0.4220762 0.4223390 0.009584546
## 3 0.01000000      3 0.4000000 0.4113009 0.009484175
##
## Variable importance
##          P90        zP90      median prop_of_years      prop_abs
##          22           18          16          15          14
##          zmedian
##          14
##
## Node number 1: 9323 observations,    complexity param=0.5779238
## predicted class= good expected loss=0.4081304  P(node) =1
##   class counts: 5518 3805
##   probabilities: 0.592 0.408
##   left son=2 (5466 obs) right son=3 (3857 obs)
## Primary splits:
##          P90        < 3.4      to the right, improve=1869.391, (0 missing)
##          zP90        < 2.35     to the right, improve=1833.206, (0 missing)
##          prop_of_years < 0.3226179 to the right, improve=1613.383, (0 missing)
##          median        < 1.75     to the right, improve=1511.146, (0 missing)
##          zmedian        < 0.75     to the right, improve=1425.147, (0 missing)
## Surrogate splits:
##          zP90        < 2.25     to the right, agree=0.921, adj=0.809, (0 split)
##          median        < 1.75     to the right, agree=0.887, adj=0.728, (0 split)
##          prop_of_years < 0.3689931 to the right, agree=0.863, adj=0.670, (0 split)
##          zmedian        < 0.25     to the right, agree=0.847, adj=0.631, (0 split)
##          prop_abs        < 0.9980539 to the left,  agree=0.834, adj=0.599, (0 split)
##
## Node number 2: 5466 observations,    complexity param=0.01103811
## predicted class= good expected loss=0.1421515  P(node) =0.586292
##   class counts: 4689 777
##   probabilities: 0.858 0.142
##   left son=4 (3043 obs) right son=5 (2423 obs)
```

```

## Primary splits:
##      P90          < 8.65      to the right, improve=127.77870, (0 missing)
##      zP90         < 5.35      to the right, improve=126.15280, (0 missing)
##      prop_list_one < 0.0003182686 to the right, improve=100.01500, (0 missing)
##      median        < 2.75      to the right, improve= 97.94465, (0 missing)
##      zmedian        < 2.75      to the right, improve= 93.53082, (0 missing)
## Surrogate splits:
##      zP90          < 8.35      to the right, agree=0.953, adj=0.894, (0 split)
##      median         < 3.75      to the right, agree=0.887, adj=0.746, (0 split)
##      zmedian        < 2.75      to the right, agree=0.858, adj=0.681, (0 split)
##      prop_of_years  < 0.7473404 to the right, agree=0.803, adj=0.555, (0 split)
##      prop_abs       < 0.9921345 to the left,  agree=0.758, adj=0.454, (0 split)
##
## Node number 3: 3857 observations
##   predicted class=bad    expected loss=0.2149339  P(node) =0.413708
##   class counts:  829  3028
##   probabilities: 0.215 0.785
##
## Node number 4: 3043 observations
##   predicted class= good  expected loss=0.04567861  P(node) =0.3263971
##   class counts:  2904  139
##   probabilities: 0.954 0.046
##
## Node number 5: 2423 observations,   complexity param=0.01103811
##   predicted class= good  expected loss=0.2633099  P(node) =0.2598949
##   class counts:  1785  638
##   probabilities: 0.737 0.263
##   left son=10 (2057 obs) right son=11 (366 obs)
## Primary splits:
##      prop_abs        < 0.9866517      to the right, improve=106.49830, (0 missing)
##      prop_list_one   < 0.02390476     to the right, improve= 55.40615, (0 missing)
##      zP90            < 1.05          to the right, improve= 30.12334, (0 missing)
##      prop_of_years   < 0.1348485     to the right, improve= 27.65530, (0 missing)
##      P90             < 4.85          to the right, improve= 23.65390, (0 missing)
## Surrogate splits:
##      prop_of_years  < 0.9678538     to the left,   agree=0.851, adj=0.016, (0 split)
##      zmedian         < 4.25          to the left,   agree=0.850, adj=0.008, (0 split)
##      zP90            < 8.15          to the left,   agree=0.850, adj=0.008, (0 split)
##
## Node number 10: 2057 observations
##   predicted class= good  expected loss=0.2007778  P(node) =0.2206371
##   class counts:  1644  413
##   probabilities: 0.799 0.201
##
## Node number 11: 366 observations
##   predicted class=bad    expected loss=0.3852459  P(node) =0.03925775
##   class counts:  141   225
##   probabilities: 0.385 0.615

# plot tree
plot(fit, uniform=TRUE,
      main = "'Rules of thumb' for species occupancy modelling",
      margin = .1)
text(fit, use.n = TRUE,

```

```
all = TRUE,
cex = .7)
```

'Rules of thumb' for species occupancy modelling

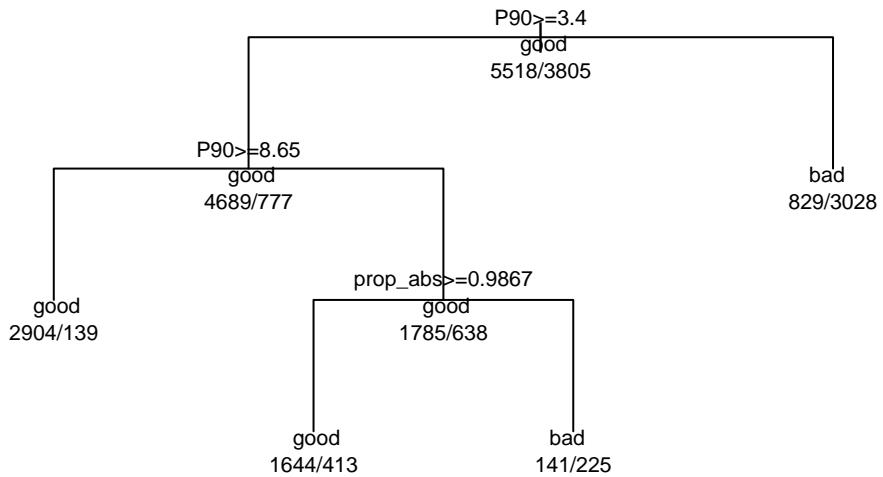


Figure 12: Unpruned decision tree for species occupancy modelling

```
# prune the tree
pfit <- prune(fit, cp = fit$cptable[which.min(fit$cptable[, "xerror"]), "CP"])

# plot the pruned tree
plot(pfit, uniform=TRUE,
      main="Pruned 'Rules of thumb' for species occupancy modelling",
      margin = .1)
text(pfit, use.n = TRUE, all = TRUE, cex = .7)
```

Pruned 'Rules of thumb' for species occupancy modelling

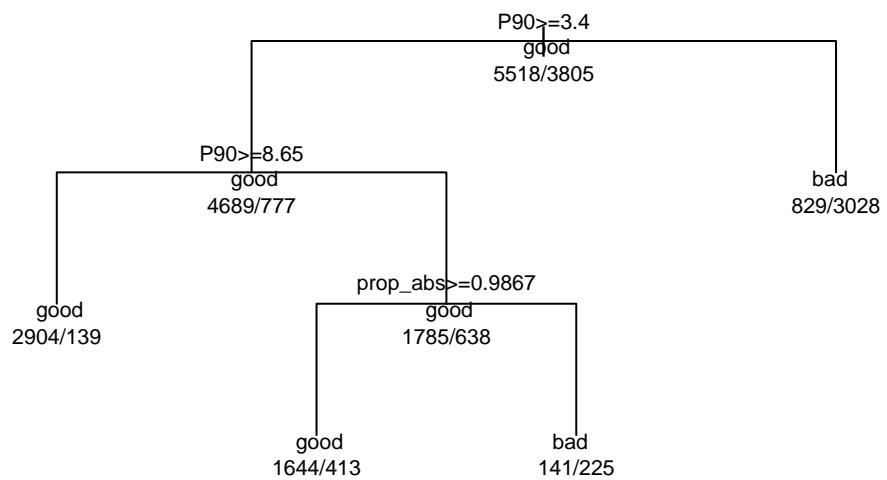


Figure 13: Pruned decision tree for species occupancy modelling

The decision tree with equal weighting to good and bad results shows a recall (chance of correct prediction) of 82% for good models and 85% for bad models, when the good/bad classifications are equally weighted. The precision (chance of classification being correct) for good models is 89% and 78% for bad models.

To explain this decision tree in words:

After removing all data from sites which never saw a repeat visit in subsequent year for any species within the taxonomic group (see *Input Data Requirement*):

- $P90 >= 3.4$: The 90th percentile of number of visits per year *in which there was at least one visit* must be greater than 3.4 or the model will definitely be classified as bad
 - For example, a species with 11 records, for which the second largest record was 4, would meet this target as the P90 would be 4.
 - A species with 2 years with records, with 1 record in one year and 4 in the other would have a P90 of 3.7, so would meet this target.
 - A species with 1 year with 4 records would have a P90 of 4, so would also meet this requirement.

If the data meets this criterion, it still must meet one of these two criteria to be classified as good:

- $P90 >= 8.65$ OR
- $prop_abs >= 0.9867$. Prop_abs is the proportion of records for the taxonomic group which do not include the species of interest.
 - For example, if there were 2000 records within a group, and 40 records of the species of interest, prop_abs would be $(2000-40)/2000 = .98$. This species would therefore not meet this criterion. However, a species with 40 records within a group of 4000 records would meet this criterion ($prop_abs = .99$).

Lets assess the quality of this decision tree

```
td <- trendsData

td$Node1 <- td$P90>=3.4
td$Node2 <- td$P90>=8.65
td$Node3 <- td$prop_abs>=0.9867

td$data_good <- rep('bad',length(td$bad))
td$data_good[(td$Node1)&(td$Node2|td$Node3)] <- 'good'

td$model_good <- rep('good',length(td$bad))
td$model_good[td$bad=='bad'] <- 'bad'
```

Below are two pie charts showing how many of our 9323 species datasets have been classified as likely to produce ‘good’ or ‘bad’ models by this decision tree, and what percentage of those models were actually good (above the average precision and convergence thresholds) or bad (below the thresholds).

Classifying the data as likely to produce a good model:

$P90 >= 3.4 \& (P90 >= 8.65 \text{ OR } prop_abs >= 0.9867)$

```
plot_pie(td[(td$data_good)=='good',],
         nrow(td),
         "Classified as good")
```

Classifying the data as likely to produce a ‘bad’ model:

$P90 < 3.4 \text{ OR } (P90 >= 3.4 \& P90 < 8.65 \& prop_abs < 0.9867)$

```
plot_pie(td[(td$data_good)=='bad',],
         nrow(td),
         "Classified as bad")
```

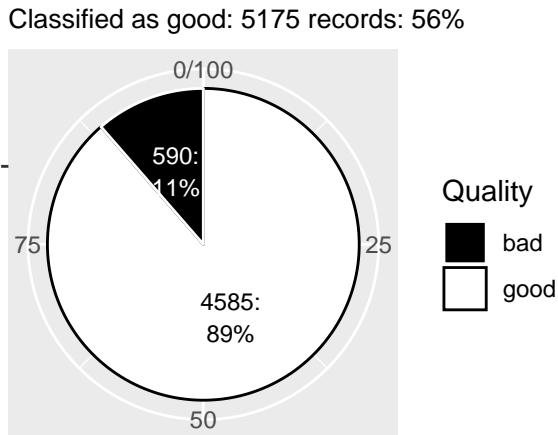


Figure 14: Pie chart showing all the models which are classified as good, and what proportion of them are actually good. This shows that the precision of this classification is 89%

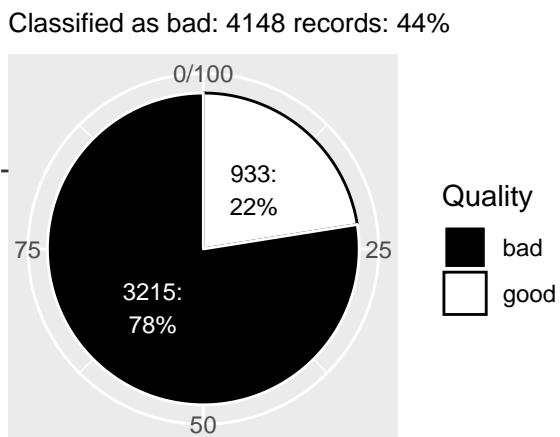


Figure 15: Pie chart showing all the models which are classified as bad, and what proportion of them are actually bad. This shows that the precision of this classification is 78%

10.1 Good:bad tree - Aspirational Target

The decision tree is computed again twice more. In these models, the penalty for misclassifying good models as bad, or vice versa, is varied.

The first one below is with a 10:1 good:bad ratio i.e. it is 10x more important to correctly classify good models as good. This means that any models which are classified as good are very likely to be good. As such, this decision tree can be considered an aspirational target.

```
# produce a decision tree for goodweight = 10
badweight <- 1
goodweight <- 10
fit_good <- rpart(bad ~ median + P90 + zmedian + zP90 + visits_median + visits_P70 +
                    prop_list_one + prop_abs,
                    method = "class",
                    data = trendsData,
                    parms = list(loss = matrix(c(0,goodweight,badweight,0),ncol = 2)))

# prune the tree
pfit_good <- prune(fit_good,
                     cp = fit_good$cptable[which.min(fit_good$cptable[, "xerror"]),"CP"])

# plot the pruned tree
plot(pfit_good, uniform=TRUE,
      main="'Rules of thumb' for species occupancy modelling, 10:1 good:bad",
      margin = .1)
text(pfit_good, use.n = TRUE, all = TRUE, cex = .7)
```

'Rules of thumb' for species occupancy modelling, 10:1 good:bad

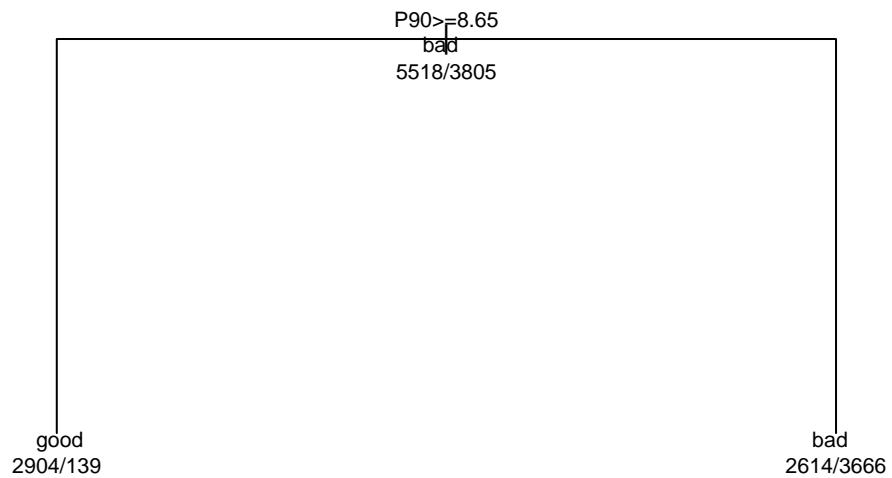


Figure 16: Decision tree for aspirational targets

In the case of the 10:1 good:bad tree, the tree is very simple. If $P90 \geq 8.65$, the data are likely to produce a good model.

The precision of this decision is 95%, while the precision of the bad classification is 58%. The recall for good models is 53% and 96% for bad models. This means only 53% of data that might produce a good model are deemed good enough. However, if data meet this threshold, they are very likely going to produce a good model.

To show how this split applies to all the data, a plot is shown below showing how this split applies to all species across all taxonomic groups in the database.

As can be seen, this split does a good job of removing most bad models, with the exception of some models from taxonomic groups which in general do not have many records per species, such as FungusGnats and RoveBeetles.

```
td_taxa <- trendsData[as.character(trendsData$Taxa)==  
                      as.character(trendsData$Taxa_Root),]  
td_taxa <- td_taxa[td_taxa$P90!=0,]  
  
## A little bit of data manipulation is required to make a nice graph  
td_taxa$colour <- rep('blue',length(td_taxa$bad))  
td_taxa$colour[td_taxa$bad=='bad'] <- 'red'  
td_taxa$cat <- paste0(td_taxa$Taxa_Root, ' - ',td_taxa$bad)  
td_taxa$alpha <- 0.2  
taxa_list <- as.character(unique(td_taxa$Taxa))  
taxa_list <- taxa_list[!is.na(taxa_list)]  
tmp_bind <- td_taxa[1:length(taxa_list),]  
tmp_bind$P90 <- 1  
tmp_bind$cat <- paste0(sort(taxa_list), ' ' )  
tmp_bind$alpha <- 0  
td_taxa <- rbind(td_taxa,tmp_bind)  
names <- sort(unique(as.character(td_taxa$Taxa_Root)))  
names_long <- c()  
for(i in 1:length(names)){  
  names_long[i*3-2] <- ''  
  names_long[i*3-1] <- names[i]  
  names_long[i*3] <- ''  
}  
}
```

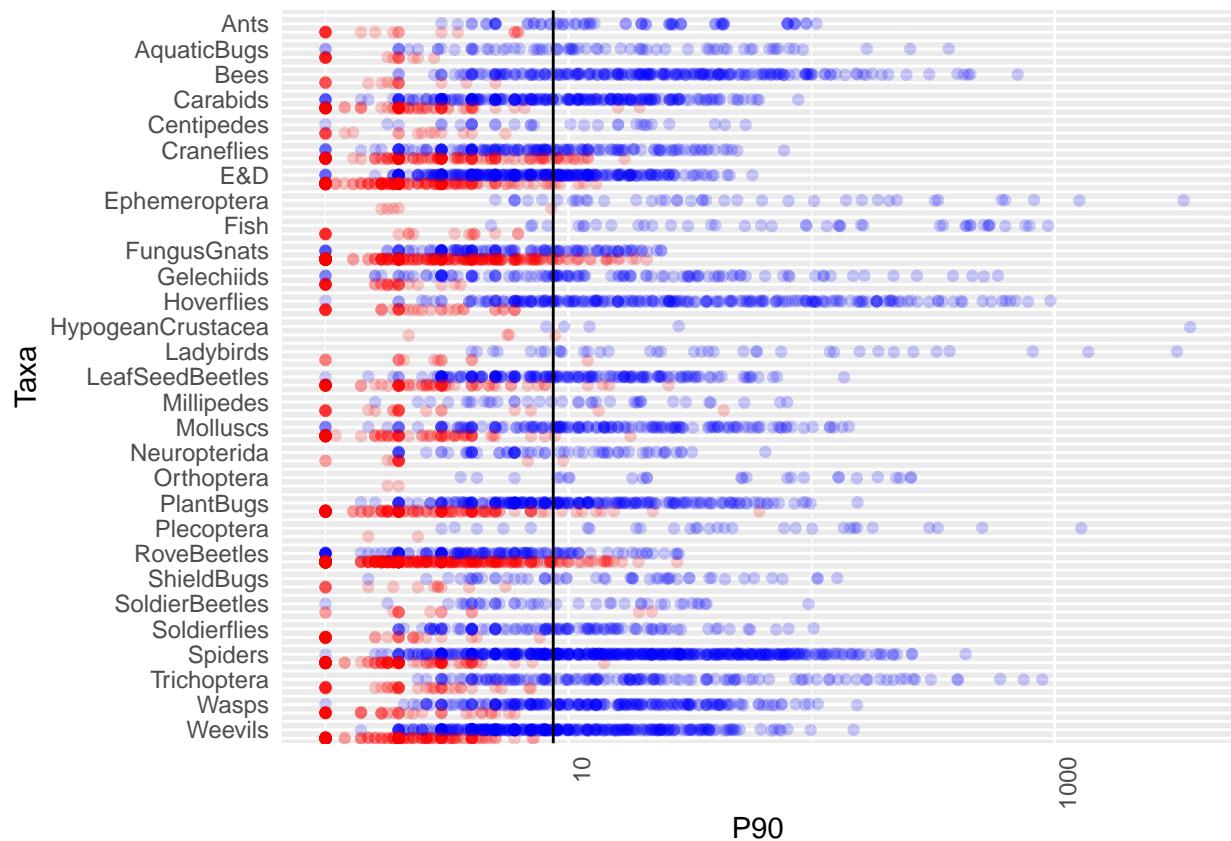


Figure 17: Graph showing all species in the database split by taxonomic group. On the y-axis are all the taxonomic groups, while the x-axis shows the P90 with the 8.65 cutoff shown as a vertical bar. Red circles represent models which have been classified as bad, while blue are good. Any circles to the right of the line represent species data which would be expected to produce a good model (95% of such datasets did in fact produce a good model). Any circles to the left represent species data which may or may not produce a good model (42% of these datasets actually produced a good model).

1:10 Good:bad tree - bare minimum to try running a model

The decision tree below is with a 1:10 good:bad ratio i.e. it is 10x more important to correctly classify bad models as bad. This means that any models which are classified as bad are very likely to be bad, while a model classified as good may or may not be actually good. As such, this decision tree can be considered a bare minimum threshold.

```
# produce a decision tree for badweight = 10
goodweight <- 1
badweight <- 10
fit_bad <- rpart(bad ~ median + P90 + zmedian + zP90 +
                  visits_median + visits_P70 + visits_P80 + visits_P90 +
                  prop_of_years + prop_repeats_grp + prop_list_one +
                  prop_abs,
                  method = "class",
                  data = trendsData,
                  parms = list(loss = matrix(c(0,goodweight,badweight,0),ncol = 2)))

# prune the tree
pfit_bad <- prune(fit_bad,
                   cp = fit_bad$cptable[which.min(fit_bad$cptable[, "xerror"]),
                                         "CP"])

# plot the pruned tree
plot(pfit_bad, uniform=TRUE,
      main="Rules of thumb' for species occupancy modelling, 1:10 good:bad",
      margin = .1)
text(pfit_bad, use.n = TRUE, all = TRUE, cex = .7)
```

'Rules of thumb' for species occupancy modelling, 1:10 good:bad

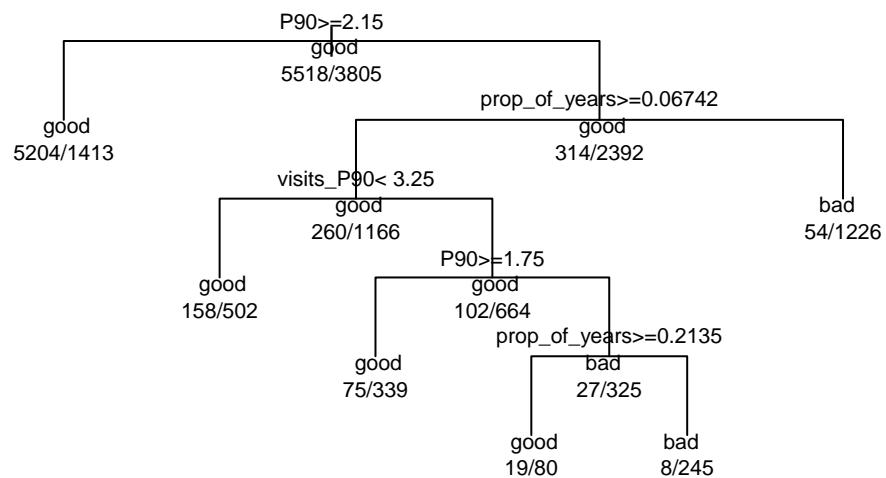


Figure 18: Decision tree for aspirational targets

The 1:10 good:bad tree is more complex, and requires some disentangling.

A dataset will be classified as good enough if it meets any of the following criteria:

- $P90 >= 2.15$
- $prop_of_years >= 0.06742$ and either:
 - $visits_P90 < 3.15$ and $P90 < 2.15$ OR
 - $2.15 > P90 >= 1.75$ and $prop_abs >= 0.9961$
- $P90 < 1.75$ and $prop_of_years >= 0.2135$

The decision tree produces a precision and recall for bad models of 96% and 40% respectively, and 71% and 99% for good models. This means that if the data does not meet these criteria, there is a 96% of it producing a poor model.

The first node alone (is $P90 >= 2.15$) gives a precision and recall for bad models of 88% and 63%, and for good models 79% and 94%. This means that if the data does not meet this threshold, there is a 88% of it producing a poor model.

If Spnvisits (number of records) is allowed back into the decision tree, the only question that needs to be asked is: are there at least 2 records? This results in a precision and recall for bad models of 97% and 21%, and 67% and 99.5% for good models, a similar classification rate for a much simpler decision tree.

How many species can we model?

For the UK

Based on the work above, we can estimate how many species we can model for taxa across the whole of the UK and across multiple habitat and regional areas.

To provide a benchmark, below are two plots for all taxa for all UK data. The first graphs shows absolute values, while the second shows the proportion of species for each taxonomic group.

```
# Extract raw data metrics, as we want to look at all species, not just those that we had matching post
RM <- read.csv('Results/metrics/ALL_rawMetrics.csv')
# Remove all last 10 yr data
RM$Taxa <- as.character(RM$Taxa)
RM <- RM[RM$Taxa==as.character(RM$Taxa_Root),]

# Calculate which datasets are likely to produce good or bad models
calc_bad <- function(df){
  df$Node1 <- df$P90>=3.4
  df$Node2 <- df$P90>=8.65
  df$Node3 <- df$prop_abs>=0.9867
  df$data_good <- df$data_aspire <- rep('bad',length(df$species))
  df$data_good[(df$Node1)&(df$Node2|df$Node3)] <- 'good'
  df$data_aspire[df$Node2] <- 'good'
  if(!is.null(df$habitat)){
    df$habitat <- as.character(df$habitat)
  }
  if(!is.null(df$ragion)){
    df$region <- as.character(df$region)
  }
  if(!is.null(df$code)){
    df$code <- as.character(df$code)
  }
  return(df)
}
```

```

}

# Function to extract data from data frame
# df = a trendsData data frame
# proportional = is the graph to show proportion of species (vs absolute numbers)
# habitat = habitat of interest
# region = region of interest, either full name or NUTS code
# zeroes = are species with no records (Spnvisits==0) to be included?
# aspirational = is the cutoff the aspirational target (10:1 good:bad)? If this is
#   false, the default 1:1 decision tree is used
# long_table = is the long table for plotting required?
num_spec <- function(df, habitat=NULL, region=NULL, proportional=FALSE,
                      long_table=TRUE, zeroes=TRUE, aspirational=FALSE){
  taxa_num <- NULL
  if(!zeroes){
    df <- df[df$Spnvisits != 0,]
  }
  for(taxa in sort(unique(df$Taxa))){
    num_spec <- length(unique(df$species[df$Taxa == taxa]))
    taxa_num <- rbind(taxa_num,
                       data.frame(taxa = taxa,
                                   num_spec = num_spec))
  }

  if(!is.null(habitat)){
    df <- df[as.character(df$habitat) == habitat,]
  }
  if(!is.null(region)){
    if(grepl('UK[A-Z]', region)){
      df <- df[as.character(df$code) == region,]
    } else {
      df <- df[as.character(df$region) == region,]
    }
  }
  taxa_count <- NULL
  for(taxa in sort(unique(taxa_num$taxa))){
    if(aspirational){
      data_good <- length(df$Taxa[df$Taxa == taxa &
                                    df$data_aspire == 'good'])
      data_bad <- length(df$Taxa[df$Taxa == taxa &
                                    df$data_aspire == 'bad'])
    } else {
      data_good <- length(df$Taxa[df$Taxa == taxa &
                                    df$data_good == 'good'])
      data_bad <- length(df$Taxa[df$Taxa == taxa &
                                    df$data_good == 'bad'])
    }
    no_data <- length(df$Taxa[df$Taxa == taxa &
                               df$Spnvisits == 0])
    tmpdf <- data.frame(taxa = taxa,
                         no_data = no_data,
                         bad = data_bad - no_data,
                         good = data_good)
  }
}

```

```

    taxa_count <- rbind(taxa_count, tmpdf)
}

taxa_count <- merge(taxa_count, taxa_num)
taxa_count$no_data <- taxa_count$num_spec - taxa_count$bad - taxa_count$good

# Convert data table for turning into graphs
taxa_melt <- melt(taxa_count, id=c('num_spec','taxa'))
taxa_prop <- taxa_count
taxa_prop$bad <- taxa_prop$bad/taxa_prop$num_spec
taxa_prop$good <- taxa_prop$good/taxa_prop$num_spec
taxa_prop$no_data <- taxa_prop$no_data/taxa_prop$num_spec
taxa_prop$no_data[taxa_prop$num_spec==0] <- 1
taxa_prop$bad[taxa_prop$num_spec==0] <- 0
taxa_prop$good[taxa_prop$num_spec==0] <- 0
taxa_prop_melt <- melt(taxa_prop, id=c('num_spec','taxa'))
if(proportional){
  taxa_melt <- taxa_prop_melt
}
if(long_table){
  return(taxa_melt)
} else {
  return(taxa_count)
}
}

# Function to plot data extracted from data frame from num_spec function
# df = a data frame from num_spec
# prefix = a prefix to the title to describe the area or habitat of interest
stack_taxa <- function(df,prefix=NULL){
  if('no_data' %in% as.character(unique(df$variable))){
    colours = c('#CCCCCC','red', '#9999FF')
  } else {
    colours = c('red', '#9999FF')
  }
  if(max(df$value)<=1){
    ylabel <- 'Proportion of Species'
  } else {
    ylabel <- 'Number of Species'
  }
  if(!is.null(prefix)){
    title <- paste0(prefix,': ',ylabel,' which can be modelled')
  } else {
    title <- paste(ylabel,'which can be modelled')
  }
  stack_records(df,colours,ylabel,title)
}

RM <- calc_bad(RM)

# Plot the basic graph
num_spec(RM) %>% stack_taxa()

```

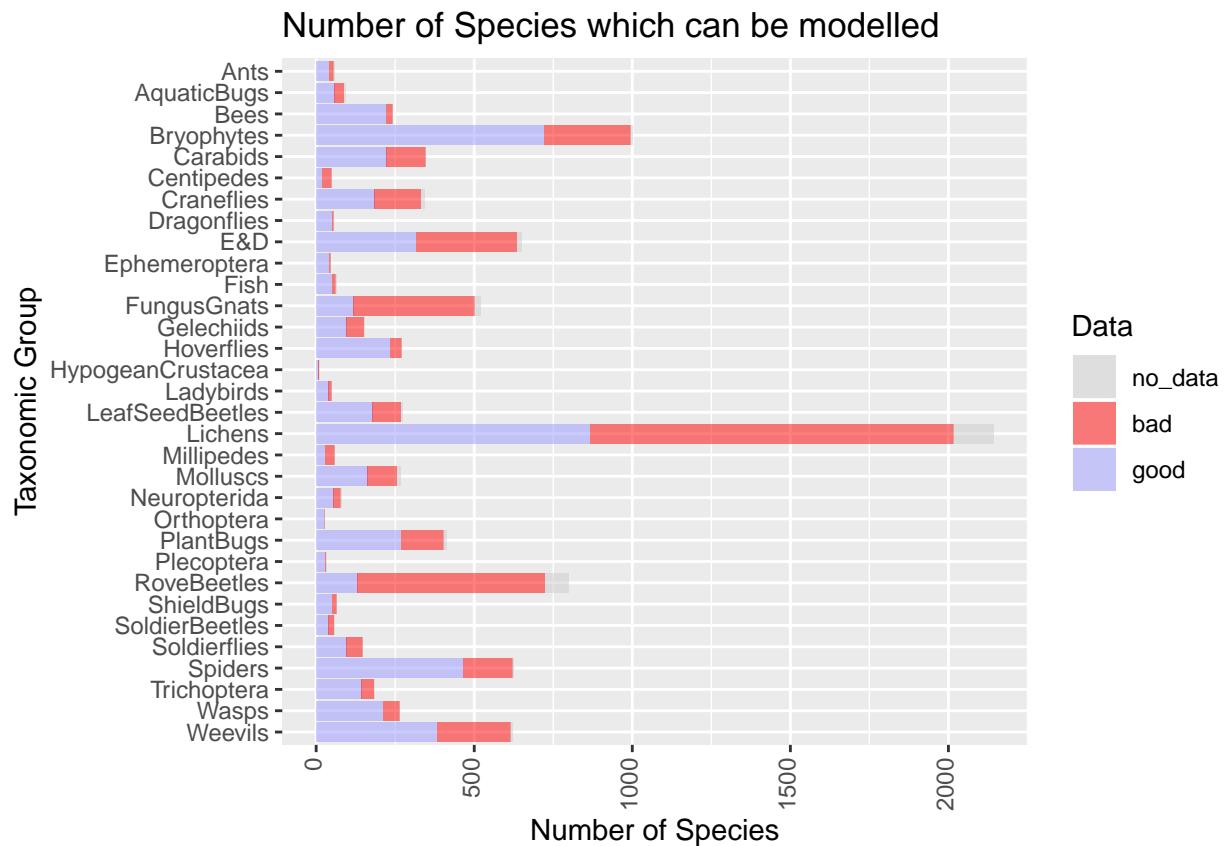


Figure 19: Plot of number of species which can be modelled, split by taxonomic group. Blue shows species which can be modelled, red those which cannot. Grey are those for which there is no data, due to there being no records which meet the requirement of coming from sites with repeat visits in any subsequent years.

```
num_spec(RM,proportional=TRUE) %>% stack_taxa()
```

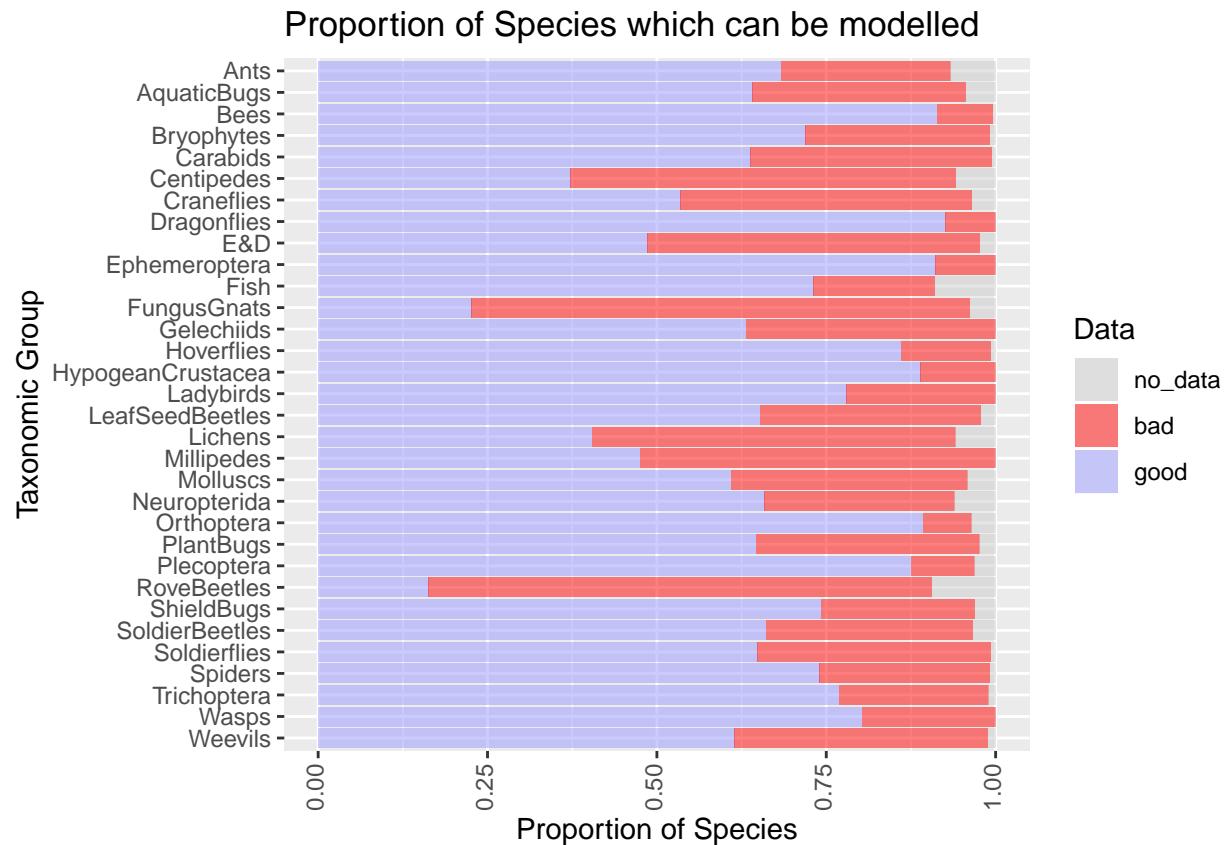


Figure 20: Plot of number of species which can be modelled, as a proportion of the total number of species for each taxonomic group.

```
# Plot aspirational graph
num_spec(RM,aspirational=TRUE,proportional=TRUE) %>%
  stack_taxa(prefix = 'Aspirational')
```

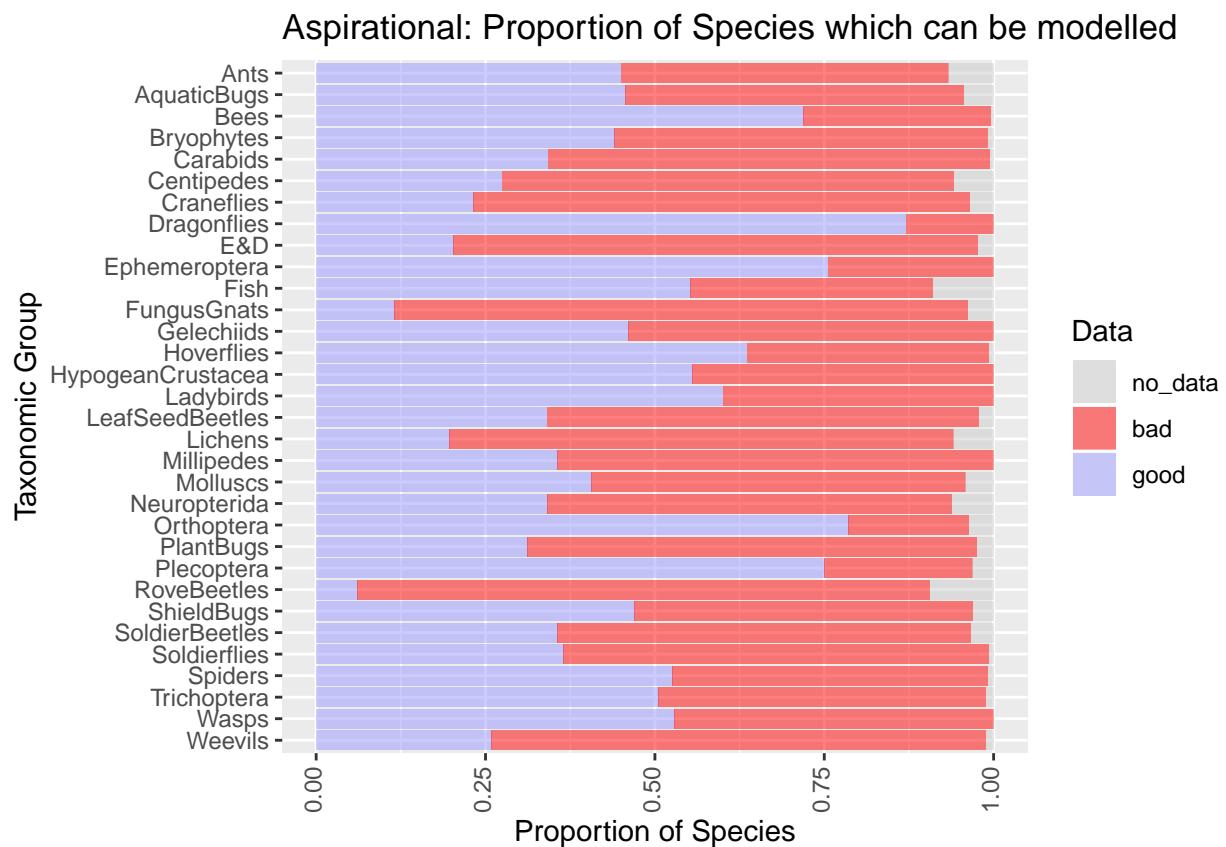


Figure 21: Plot of number of species which can be modelled using the aspirational target. As this is the aspirational target, the number of species which can be modelled is lower

Table 1: List of available habitats

Habitat	Description	Habitat_code
Broadleaf woodland	Broadleaved, Mixed and Yew Woodland	BLW
Coniferous woodland	Coniferous Woodland	CWL
Arable	Arable and Horticultural	A
Improved grassland	Improved Grassland	IG
Semi-natural grassland	Neutral Grassland, Calcareous Grassland, Acid Grassland, Fen, Marsh and Swamp	SNG
Mountain, heath, bog	Heather, Heather Grassland, Bog, Inland Rock	MHB
Coastal	Supra-littoral Rock, Supra-littoral Sediment, Littoral Rock, Littoral Sediment, Saltmarsh	C
Freshwater	Freshwater	FW
Built-up areas and gardens	Built-up Areas and Gardens	BU

By habitat

The possible habitats are shown in Table 1.

```
# Display the list of possible habitats
hab_list <- read.csv('Habitat/Habitat_List.csv')
column_spec(kable(hab_list, format = 'latex',
                  caption = 'List of available habitats'),
            2, width = "8cm")
```

To calculate the raw metrics for each habitat, the squares in the UK all need to be assigned to a habitat. To do so, a habitat datafile is read in and the range of percentage cover for each habitat is extracted. This produces a list of coverage for each habitat such as:

- Habitat: BLW; Coverage: 1%, 1%, 2%, 3% ... 95%, 98%, 99%

From this list, the median coverage for each habitat is calculated. Each square in the UK which has above this median coverage is assigned to that habitat. Therefore it is possible for a square to contain multiple habitats or not represent any habitat.

```
# Read in land habitat coverage data file
UK_land <- readRDS(file.path('./Habitat/UK_land.rds'))
UK_land_GR <- UK_land
UK_land_GR$easting = UK_land$easting + 500
UK_land_GR$northing = UK_land$northing + 500

UK_land_GR$site <- gr_num2let(easting = UK_land_GR$easting,
                                northing = UK_land_GR$northing)
UK_land_GR$site <- paste0(substr(x = UK_land_GR$site, start = 1, stop = 4),
                           substr(x = UK_land_GR$site, start = 8, stop = 9))

##coverage is in km squared
total_coverage <- colSums(UK_land[,3:ncol(UK_land)])]

##now also want number of kmsq with any cover of each habitat
UK_land_any_hab <- ifelse(UK_land[,3:ncol(UK_land)]>0, 1, 0)
total_coverage_squares <- colSums(UK_land_any_hab)

##now look at coverage for each of quantiles
quants <-
```

```

data.frame(apply(UK_land[,3:ncol(UK_land)], 2, function(x){
  quantile(x[x>0], probs = c(0.25, 0.50, 0.75))
}))

for(i in 1:length(quants)){
  for(j in 1:3){
    ##subset data in quant to quantile ranges (25%, 50%, 75%)
    sub_UK <- subset(UK_land, UK_land[,names(quants)[i]] > quants[j ,i])
    sub_UK_any_hab <- ifelse(sub_UK[,3:ncol(sub_UK)] > 0, 1, 0)
    quants_sums[j,i] <- sum(sub_UK_any_hab[,names(quants)[i]])
  }
}

```

Now calculate the metrics for each habitat

```

hab_names <- colnames(UK_land[,3:ncol(UK_land)])
for(data_file in data_files){
  formatted_data <- dataPrep(data_file)

  spp <- unique(taxa_data$CONCEPT)
  # one group has a species called '', best to remove this (single) record
  spp <- spp[spp != '']

  # Create combined dataframe
  species_obs <-
    formatted_data$occDetdata %>% left_join(formatted_data$spp_vis, by="visit")
  species_obs$site <- as.character(species_obs$site)
  # Combine with habitat data
  species_obs <- species_obs %>% left_join(UK_land_GR, by="site")
  ## Remove NA's
  species_obs <- species_obs[!is.na(species_obs$CWL),]

  yps <- rowSums(acast(species_obs, site ~ year, length, value.var = 'L') > 0)
  sites_to_include <- names(yps[yps >= 2])
  species_obs <- species_obs[(species_obs$site %in% sites_to_include),]

  spp_ch <- as.character(spp)
  raw_metrics <- data.frame()

  # At the group level, determine records with above average habitat coverage
  for(hab in hab_names){
    ## Find visits with above median coverage for this habitat
    above_av <- species_obs[,colnames(species_obs) == hab]>=
      quants[rownames(quants) == '50%', colnames(quants) == hab]
    species_obs[,paste0(hab, '_hab')] <- above_av
  }

  # For each species, find the visits with just that species recorded
  i <- 1
  for(sp in spp_ch){
    if((i %% 10) == 0){cat('  Species', i, 'of', length(spp_ch), '\n')}
    if(sp %in% names(species_obs)){
      taxa_rec <-
        species_obs[,names(species_obs) %in% c('visit', 'site', 'L', 'year',

```

```

                sp,hab_names,
                paste0(hab_names,'_hab'))]

for(hab in hab_names){
  ## Extract the metrics for each habitat
  function(sp, species_obs, basen, suffix = NULL, years = NULL)
  speciesMetric <- dataMetrics(hab = hab,
    sp = sp,
    basen = basename(data_file),
    species_obs = taxa_rec,
    habitat = TRUE)
  raw_metrics <- rbind(raw_metrics, speciesMetric)
}
}
i <- i+1
}
group_name <- regmatches(basename(data_file),
  regexp('[A-z]+(?=_)', 
  basename(data_file),
  perl=TRUE))
write.csv(raw_metrics,
  file = file.path('./Habitat/HabMetrics',
  paste0('habMetrics_',
  gsub('.rdata$', '.csv',
  basename(data_file)))),
  row.names = FALSE)
}

td <- data.frame()
file_list <- list.files(file.path('./Habitat/HabMetrics'),
  pattern = '.csv$',
  full.names = TRUE)

for(i in file_list){
  tmp <- read.csv(i)
  td <- rbind(td,tmp)
}

write.csv(x = td, file.path('./Habitat/HabMetrics/ALL_habMetrics.csv'))

```

The raw metrics for each habitat have now been calculated from the raw data. Using these data, the proportion of species which can be modelled for each habitat can be assessed.

```

# Read in the habitat raw metrics
RM_hab <-
  read.csv(file = file.path('Habitat/HabMetrics/ALL_habMetrics.csv'))

# Calculate which datasets are likely to produce good or bad models
RM_hab <- calc_bad(RM_hab)

# Plot up a sample graph for broad leaf woodland
num_spec(RM_hab,habitat='BLW',proportional=TRUE) %>%
  stack_taxa(prefix='B-L Woodland')

```

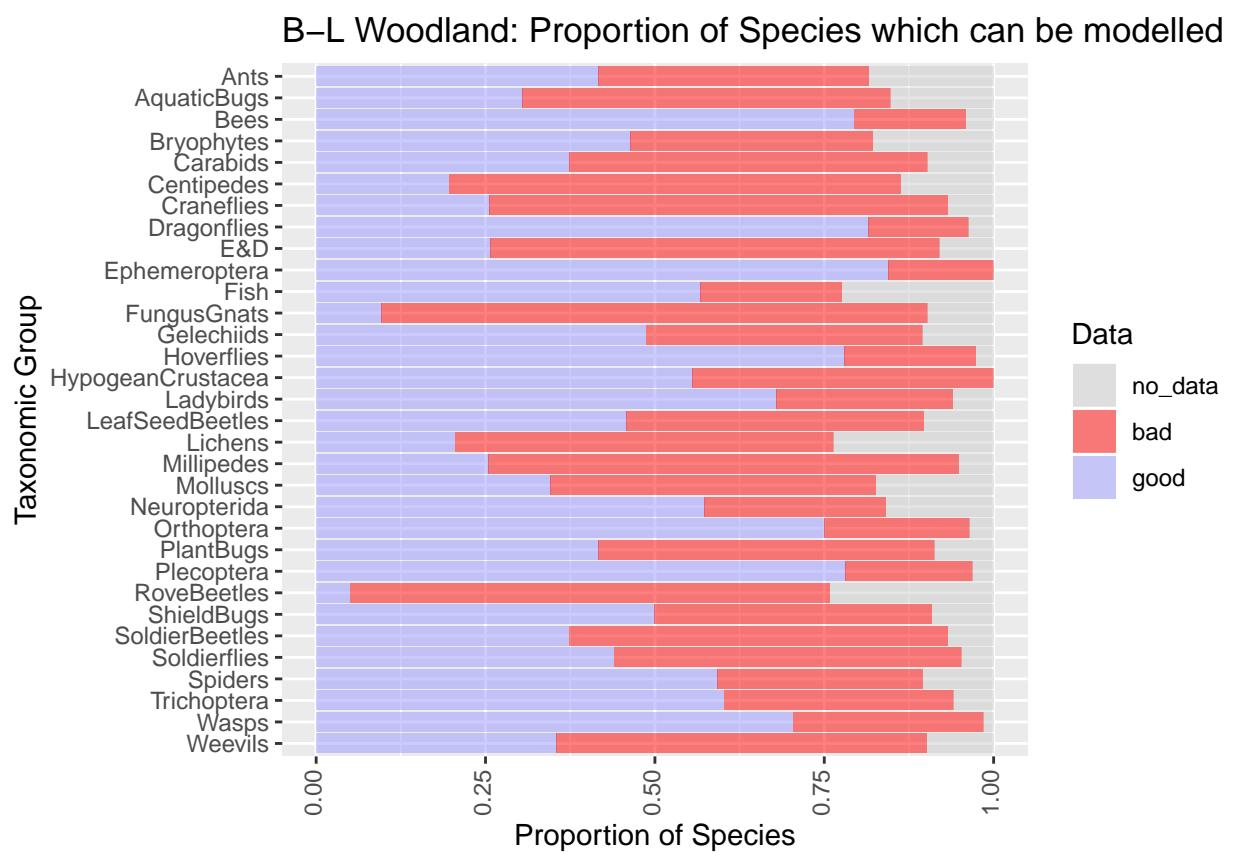


Figure 22: Proportion of species which can be modelled in broad-leaf woodland

```
# And for coastal
num_spec(RM_hab, habitat='C', proportional=TRUE) %>% stack_taxa(prefix='Coastal')
```

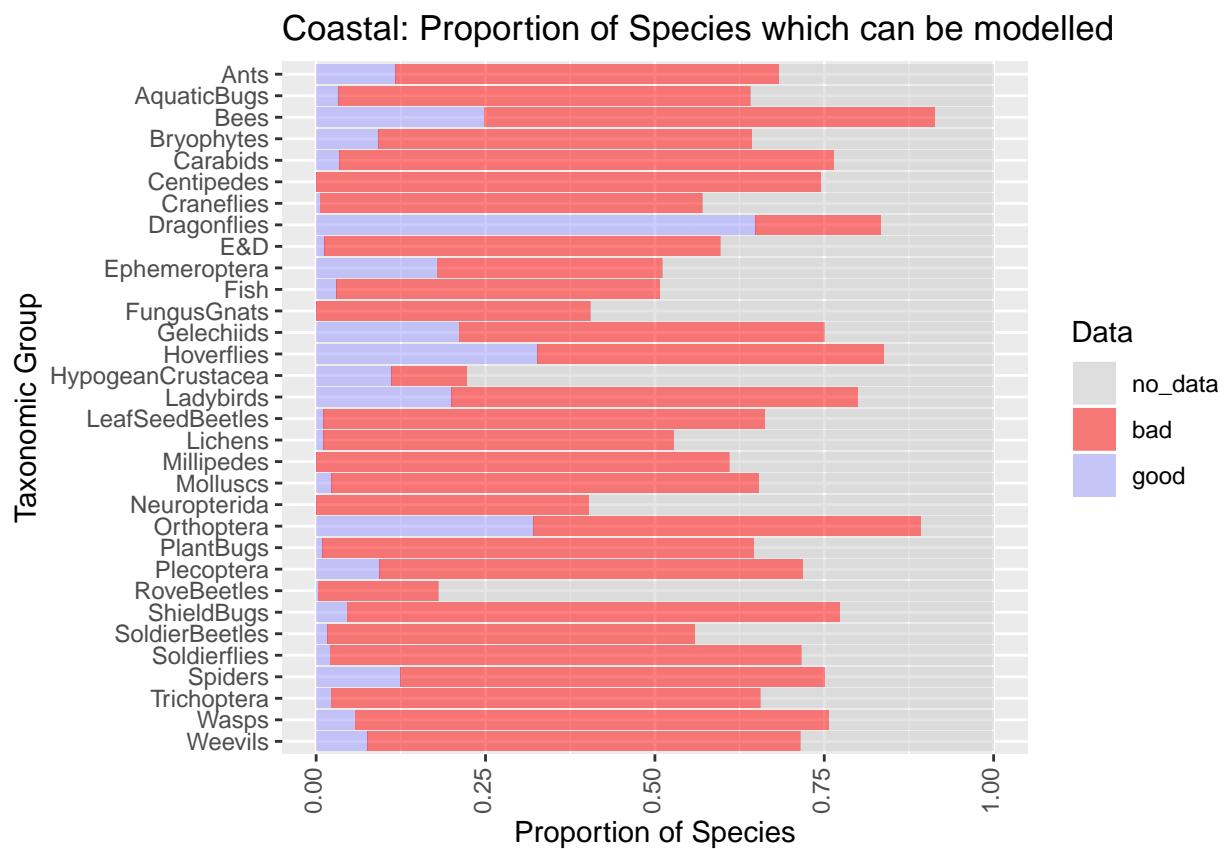


Figure 23: Proportion of species which can be modelled in coastal habitats

By region

The data can also be split by region. Splitting by region is simpler than for habitat as all squares in the UK exist within a region which can be obtained from the NUTS lookup table. The data is then subset by region and the metrics extracted for each of these regions.

```
regions <- read.csv(file.path('./Region/NUTS1_1km_dupsorted.csv'))
regions <- regions[,c(2,4)]

for(data_file in data_files){
  formatted_data <- dataPrep(data_file)
  spp <- unique(taxa_data$CONCEPT)
  spp <- spp[spp != '']

  # Create combined dataframe
  yps <- rowSums(acast(formatted_data$occDetdata,
    site ~ year, length, value.var = 'L') > 0)
  sites_to_include <- names(yps[yps >= 2])
  formatted_data$occDetdata <-
    formatted_data$occDetdata[(formatted_data$occDetdata$site %in%
      sites_to_include),]
  species_obs <-
    formatted_data$occDetdata %>% left_join(formatted_data$spp_vis,by="visit")
  species_obs$site <- as.character(species_obs$site)

  raw_metrics <- data.frame()
  species_obs <- merge(x = species_obs, y = regions, by.x= 'site', by.y = 'km_sq')

  # For each species, find the visits with just that species recorded
  for(sp in as.character(spp)){
    if(sp %in% names(species_obs)){
      taxa_rec <-
        species_obs[,names(species_obs) %in% c('visit','site','nutsname',
          'L','year',sp)]
      nutsnames <- as.character(unique(species_obs$nutsname))
      for(region in nutsnames){
        ## Extract the metrics for each habitat
        speciesMetric <- dataMetrics(hab = region,
          sp = sp,
          basen = basename(data_file),
          species_obs = taxa_rec,
          region = TRUE)
        raw_metrics <- rbind(raw_metrics, speciesMetric)
      }
    }
  }
  group_name <- regmatches(basename(data_file),
    regex('([A-z]+(?=_))', perl=TRUE))
  write.csv(raw_metrics,
    file = file.path('./Region/RegMetrics',
      paste0('regMetrics_',
        gsub('.rdata$', '.csv', basename(data_file)))),
    row.names = FALSE)
}
```

Table 2: List of regions

region	code
Northern Ireland	UKN
Scotland	UKM
North East (England)	UKC
North West (England)	UKD
Yorkshire and The Humber	UKE
Wales	UKL
West Midlands (England)	UKG
East Midlands (England)	UKF
South West (England)	UKK
South East (England)	UKJ
East of England	UKH
London	UKI

```
td <- data.frame()
file_list <- list.files(file.path('./Region/RegMetrics'),
                        full.names = TRUE)
for(i in file_list){
  td <- rbind(td,read.csv(i))
}
write.csv(x = td, file.path('./Region/RegMetrics/ALL_regMetrics.csv'))
```

Now the data has been extracted, we can read in the region data file and query the data by region. The available regions are shown in Table 2.

```
# Read in the regional raw metrics
RM_reg <-
  read.csv(file = file.path('Region/RegMetrics/ALL_regMetrics.csv'))
# Load the list of regions
region_lookup <- read.csv(file = file.path('Region/NUTS_lookup.csv'))
# Merge this lookup table with the raw data to allow it be queried by code
RM_reg <- merge(RM_reg,region_lookup)
# Calculate which datasets are likely to produce good or bad models
RM_reg <- calc_bad(RM_reg)

kable(region_lookup, caption = 'List of regions')
```

```
# Plot up a sample graph for Scotland
num_spec(RM_reg,region='Scotland',proportional=TRUE) %>% stack_taxa(prefix='Scotland')
```

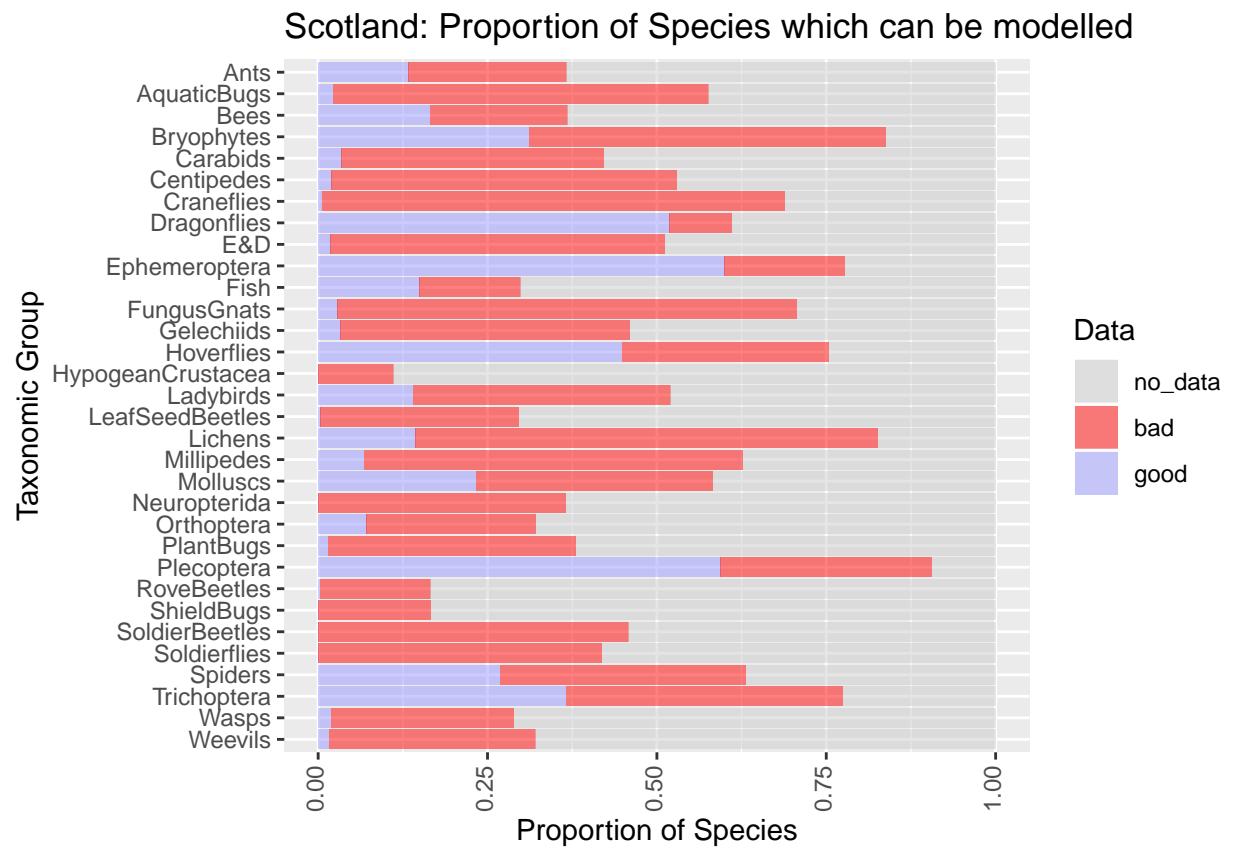


Figure 24: Proportion of species which can be modelled in Scotland

```
# And for Wales
num_spec(RM_reg,region='UKL',proportional=TRUE) %>% stack_taxa(prefix='Wales')
```

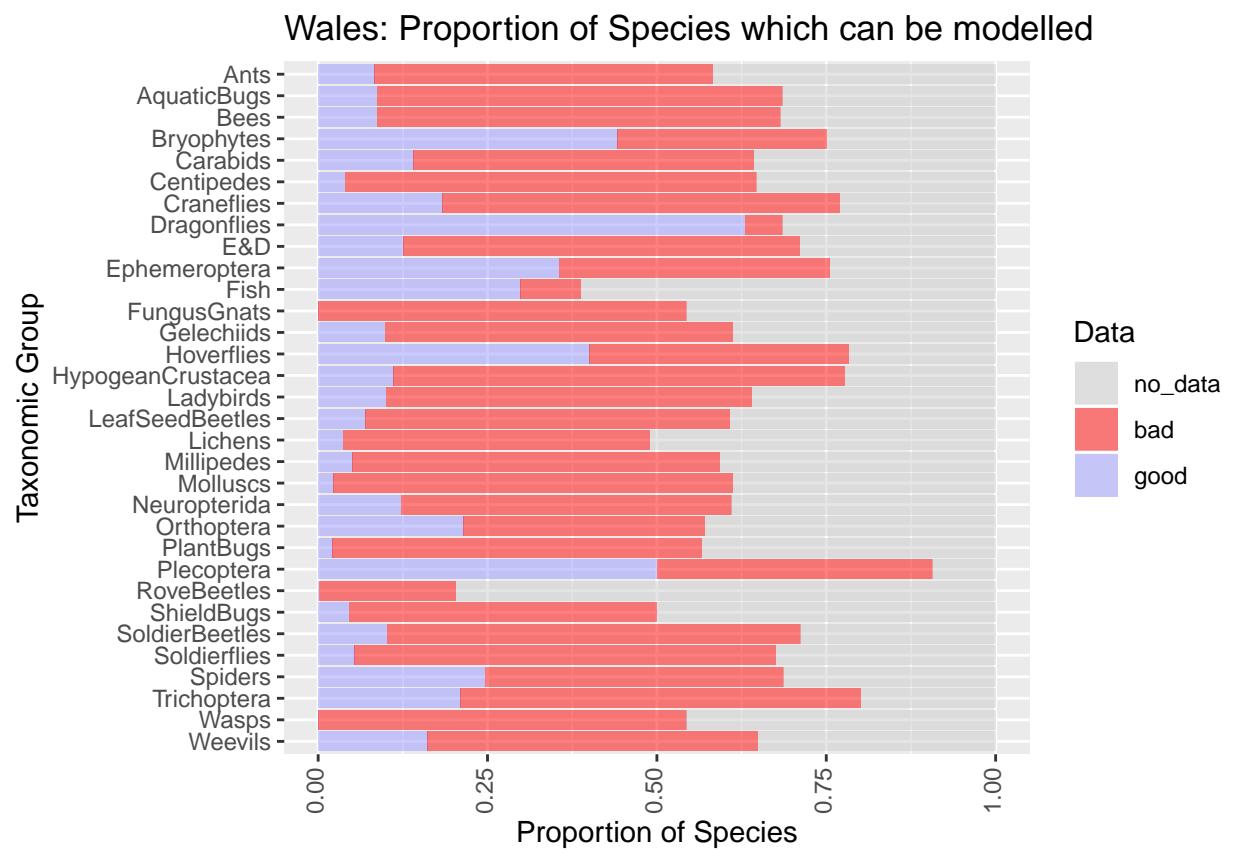


Figure 25: Proportion of species which can be modelled in Wales

Conclusions

Previously it was difficult to predict from a dataset whether it was possible to produce an occupancy model with a useful level of precision. This work has demonstrated that the criteria for making such an assessment are relatively simple and easy to apply to new datasets.

With this tool, it can be estimated how many species we can model, not just in the UK as a whole, but for any region or habitat of interest within the UK. The aspirational criteria also represent a target to aim for, to enable the modelling of any species of interest.